

Warming-up to Machine Learning, Data and Features

CS771: Introduction to Machine Learning

Piyush Rai

Plan for today

- Types of ML problems
- Typical workflow of ML problems
- Various perspectives of ML problems
- Data and Features
- Some basic operations of data and features



Keep in mind: ML is like an exam



Keep in mind: ML is like an exam

- It's the performance on the D-day which matters



Keep in mind: ML is like an exam

- It's the performance on the D-day which matters
- In an exam, our success is measured based on how well we did on the questions in the test (not on the questions we practiced on)



Keep in mind: ML is like an exam

- It's the performance on the D-day which matters
- In an exam, our success is measured based on how well we did on the questions in the test (not on the questions we practiced on)



Keep in mind: ML is like an exam

- It's the performance on the D-day which matters
- In an exam, our success is measured based on how well we did on the questions in the test (not on the questions we practiced on)
- Likewise, in ML, success of the learned model is measured based on how well it predicts/fits the future **test data** (not the training data)



Keep in mind: ML is like an exam

- It's the performance on the D-day which matters
- In an exam, our success is measured based on how well we did on the questions in the test (not on the questions we practiced on)
- Likewise, in ML, success of the learned model is measured based on how well it predicts/fits the future **test data** (not the training data)

In Machine Learning, **generalization performance** on the test data matters



Keep in mind: ML is like an exam

- It's the performance on the D-day which matters
- In an exam, our success is measured based on how well we did on the questions in the test (not on the questions we practiced on)
- Likewise, in ML, success of the learned model is measured based on how well it predicts/fits the future **test data** (not the training data)

In Machine Learning, **generalization performance** on the test data matters


Plus, of course, issues such as fairness




A Loose Taxonomy of ML



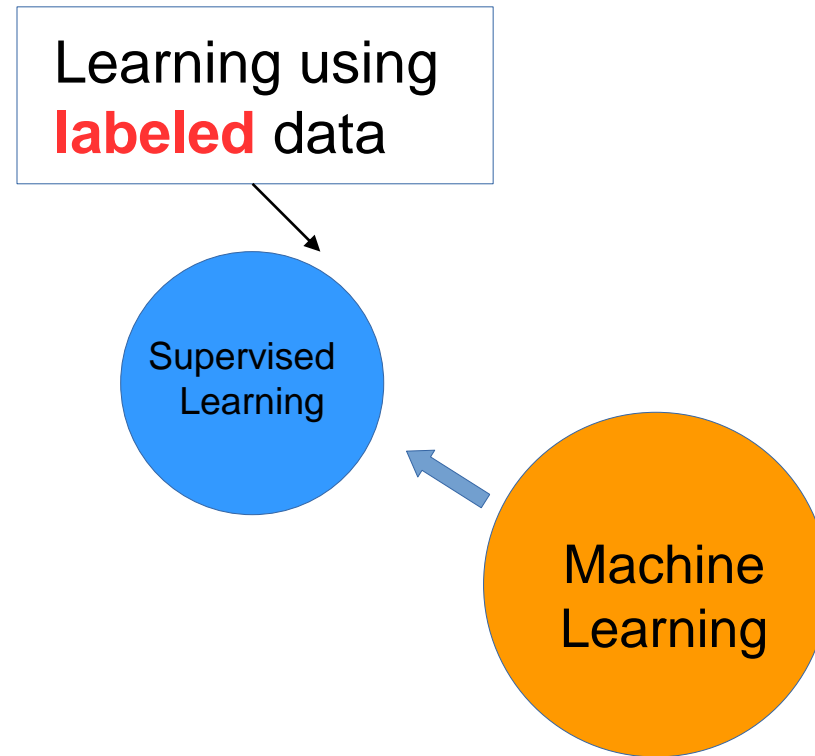
A Loose Taxonomy of ML



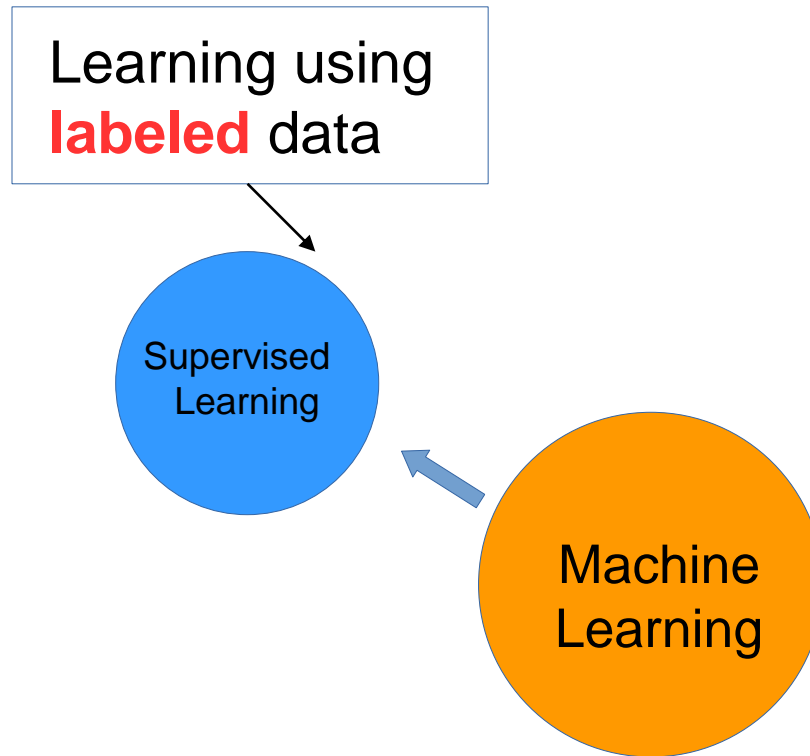
A Loose Taxonomy of ML




A Loose Taxonomy of ML



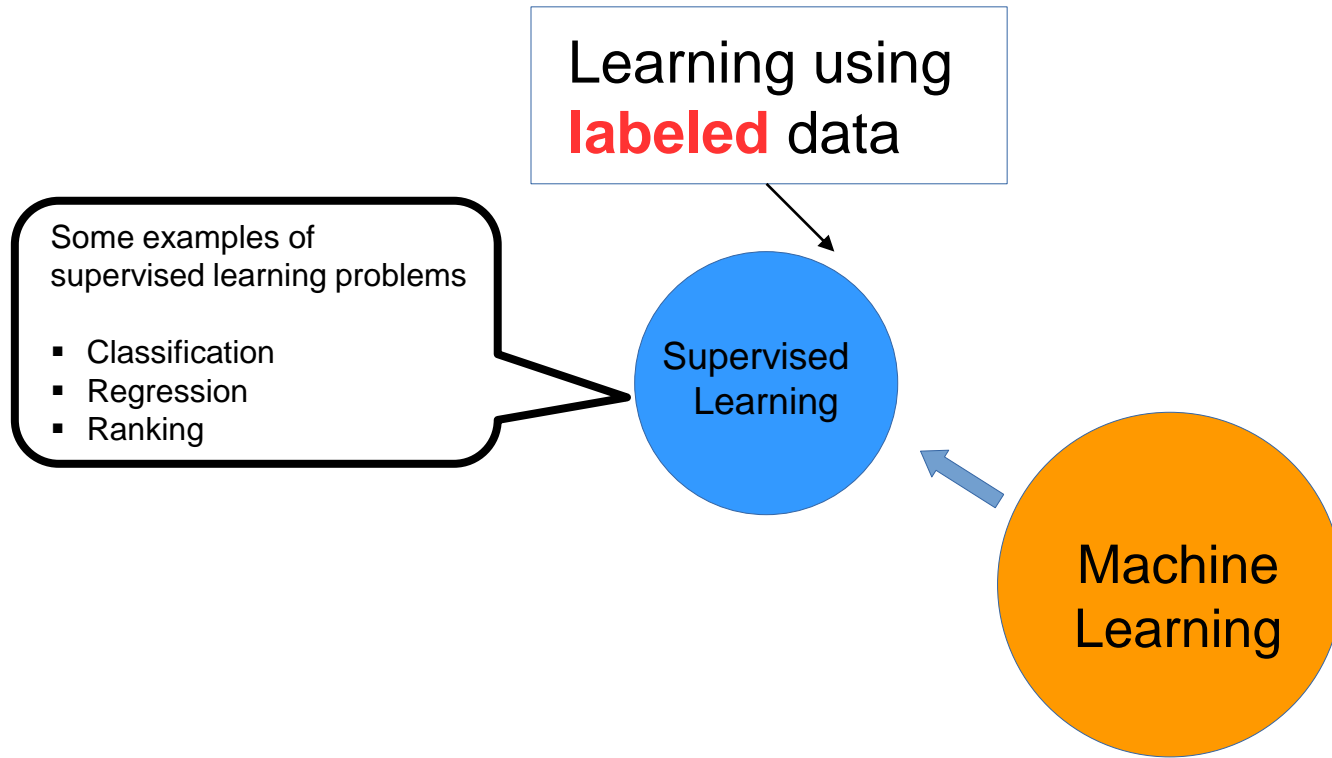
A Loose Taxonomy of ML



"Labeled" means, during training, for each input, the corresponding output is available (i.e., the machine learner is explicitly told that a cat image is of a cat)



A Loose Taxonomy of ML



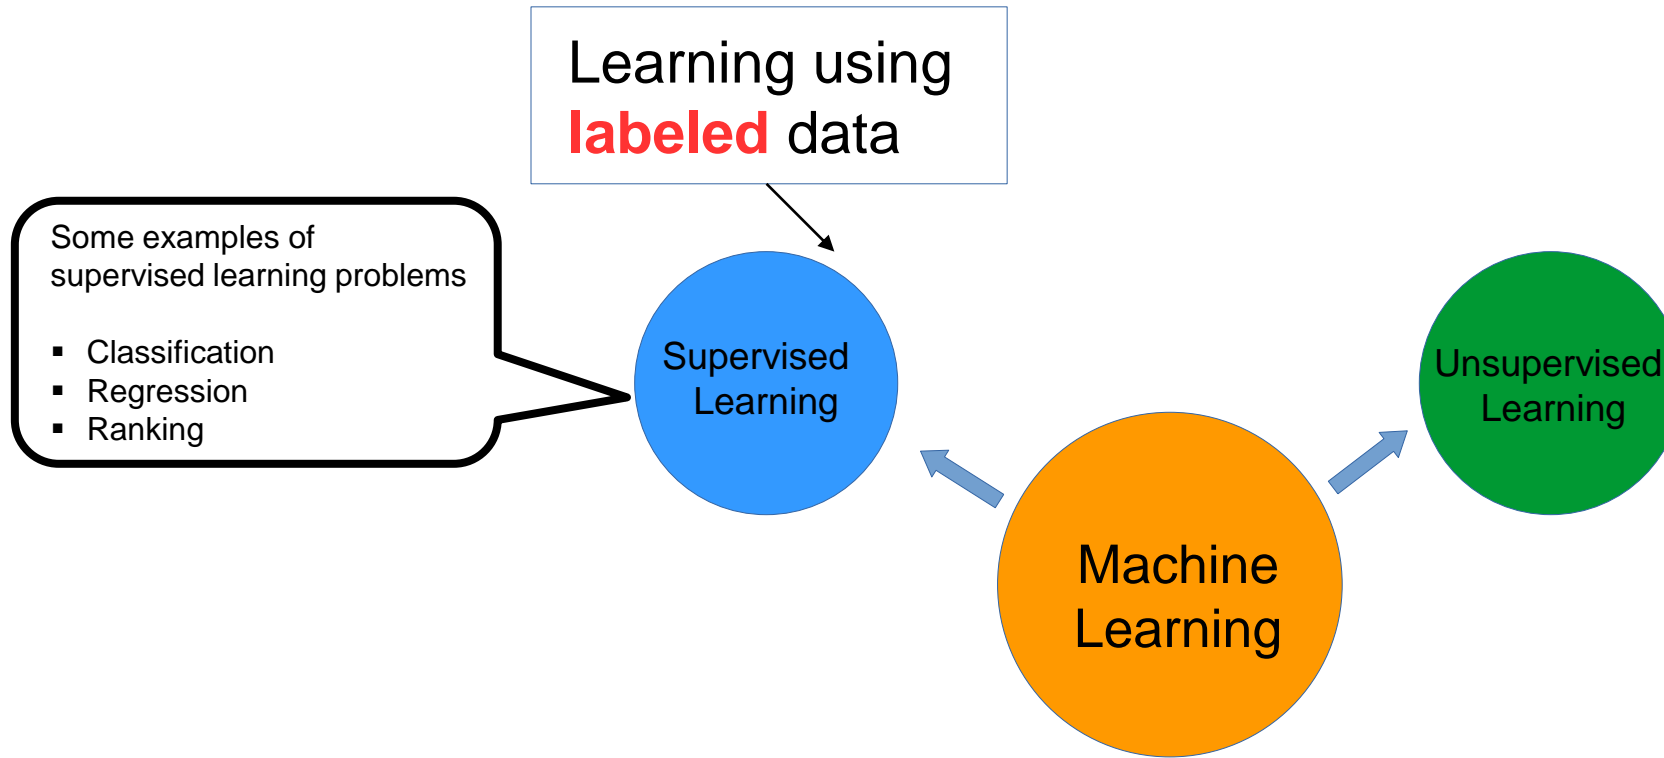
“Labeled” means, during training, for each input, the corresponding output is available (i.e., the machine learner is explicitly told that a cat image is of a cat)



4



A Loose Taxonomy of ML



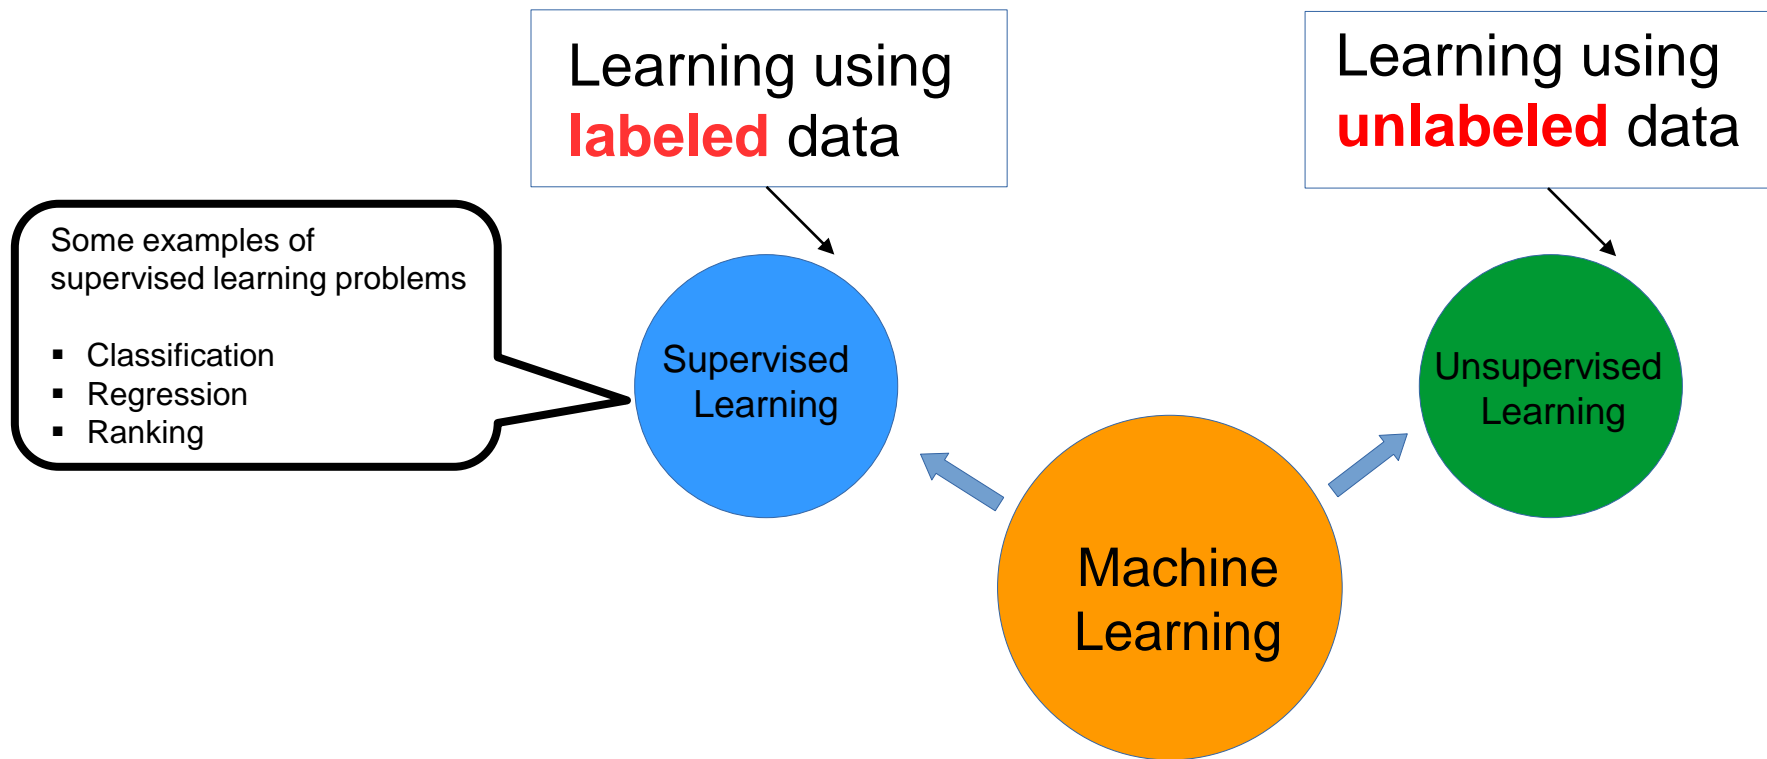
“Labeled” means, during training, for each input, the corresponding output is available (i.e., the machine learner is explicitly told that a cat image is of a cat)



4



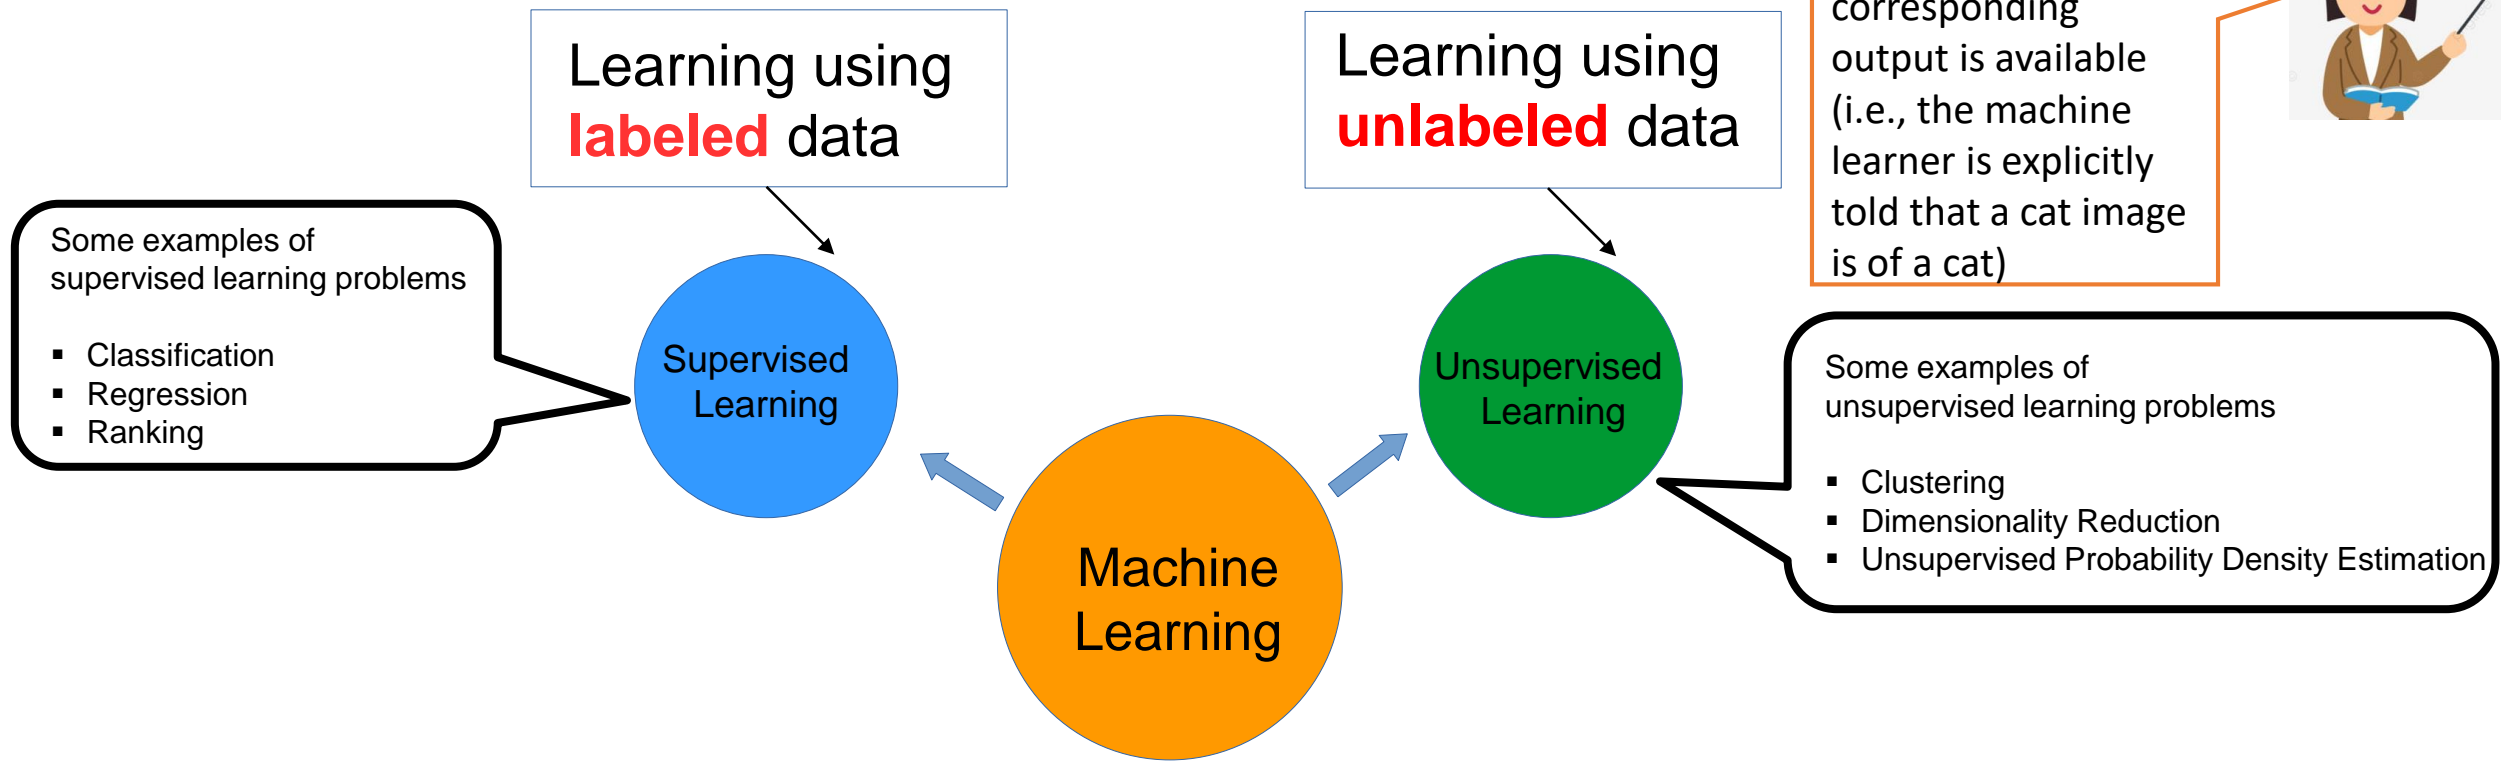
A Loose Taxonomy of ML



"Labeled" means, during training, for each input, the corresponding output is available (i.e., the machine learner is explicitly told that a cat image is of a cat)



A Loose Taxonomy of ML



A Loose Taxonomy of ML

4



“Labeled” means, during training, for each input, the corresponding output is available (i.e., the machine learner is explicitly told that a cat image is of a cat)

Learning using
labeled data

Learning using
unlabeled data

Some examples of supervised learning problems

- Classification
- Regression
- Ranking

Supervised Learning

Unsupervised Learning

Some examples of unsupervised learning problems

- Clustering
- Dimensionality Reduction
- Unsupervised Probability Density Estimation

Machine Learning

Reinforcement Learning



A Loose Taxonomy of ML

4



“Labeled” means, during training, for each input, the corresponding output is available (i.e., the machine learner is explicitly told that a cat image is of a cat)

Learning using
labeled data

Learning using
unlabeled data

Some examples of supervised learning problems

- Classification
- Regression
- Ranking

Supervised Learning

Unsupervised Learning

Some examples of unsupervised learning problems

- Clustering
- Dimensionality Reduction
- Unsupervised Probability Density Estimation

Machine Learning

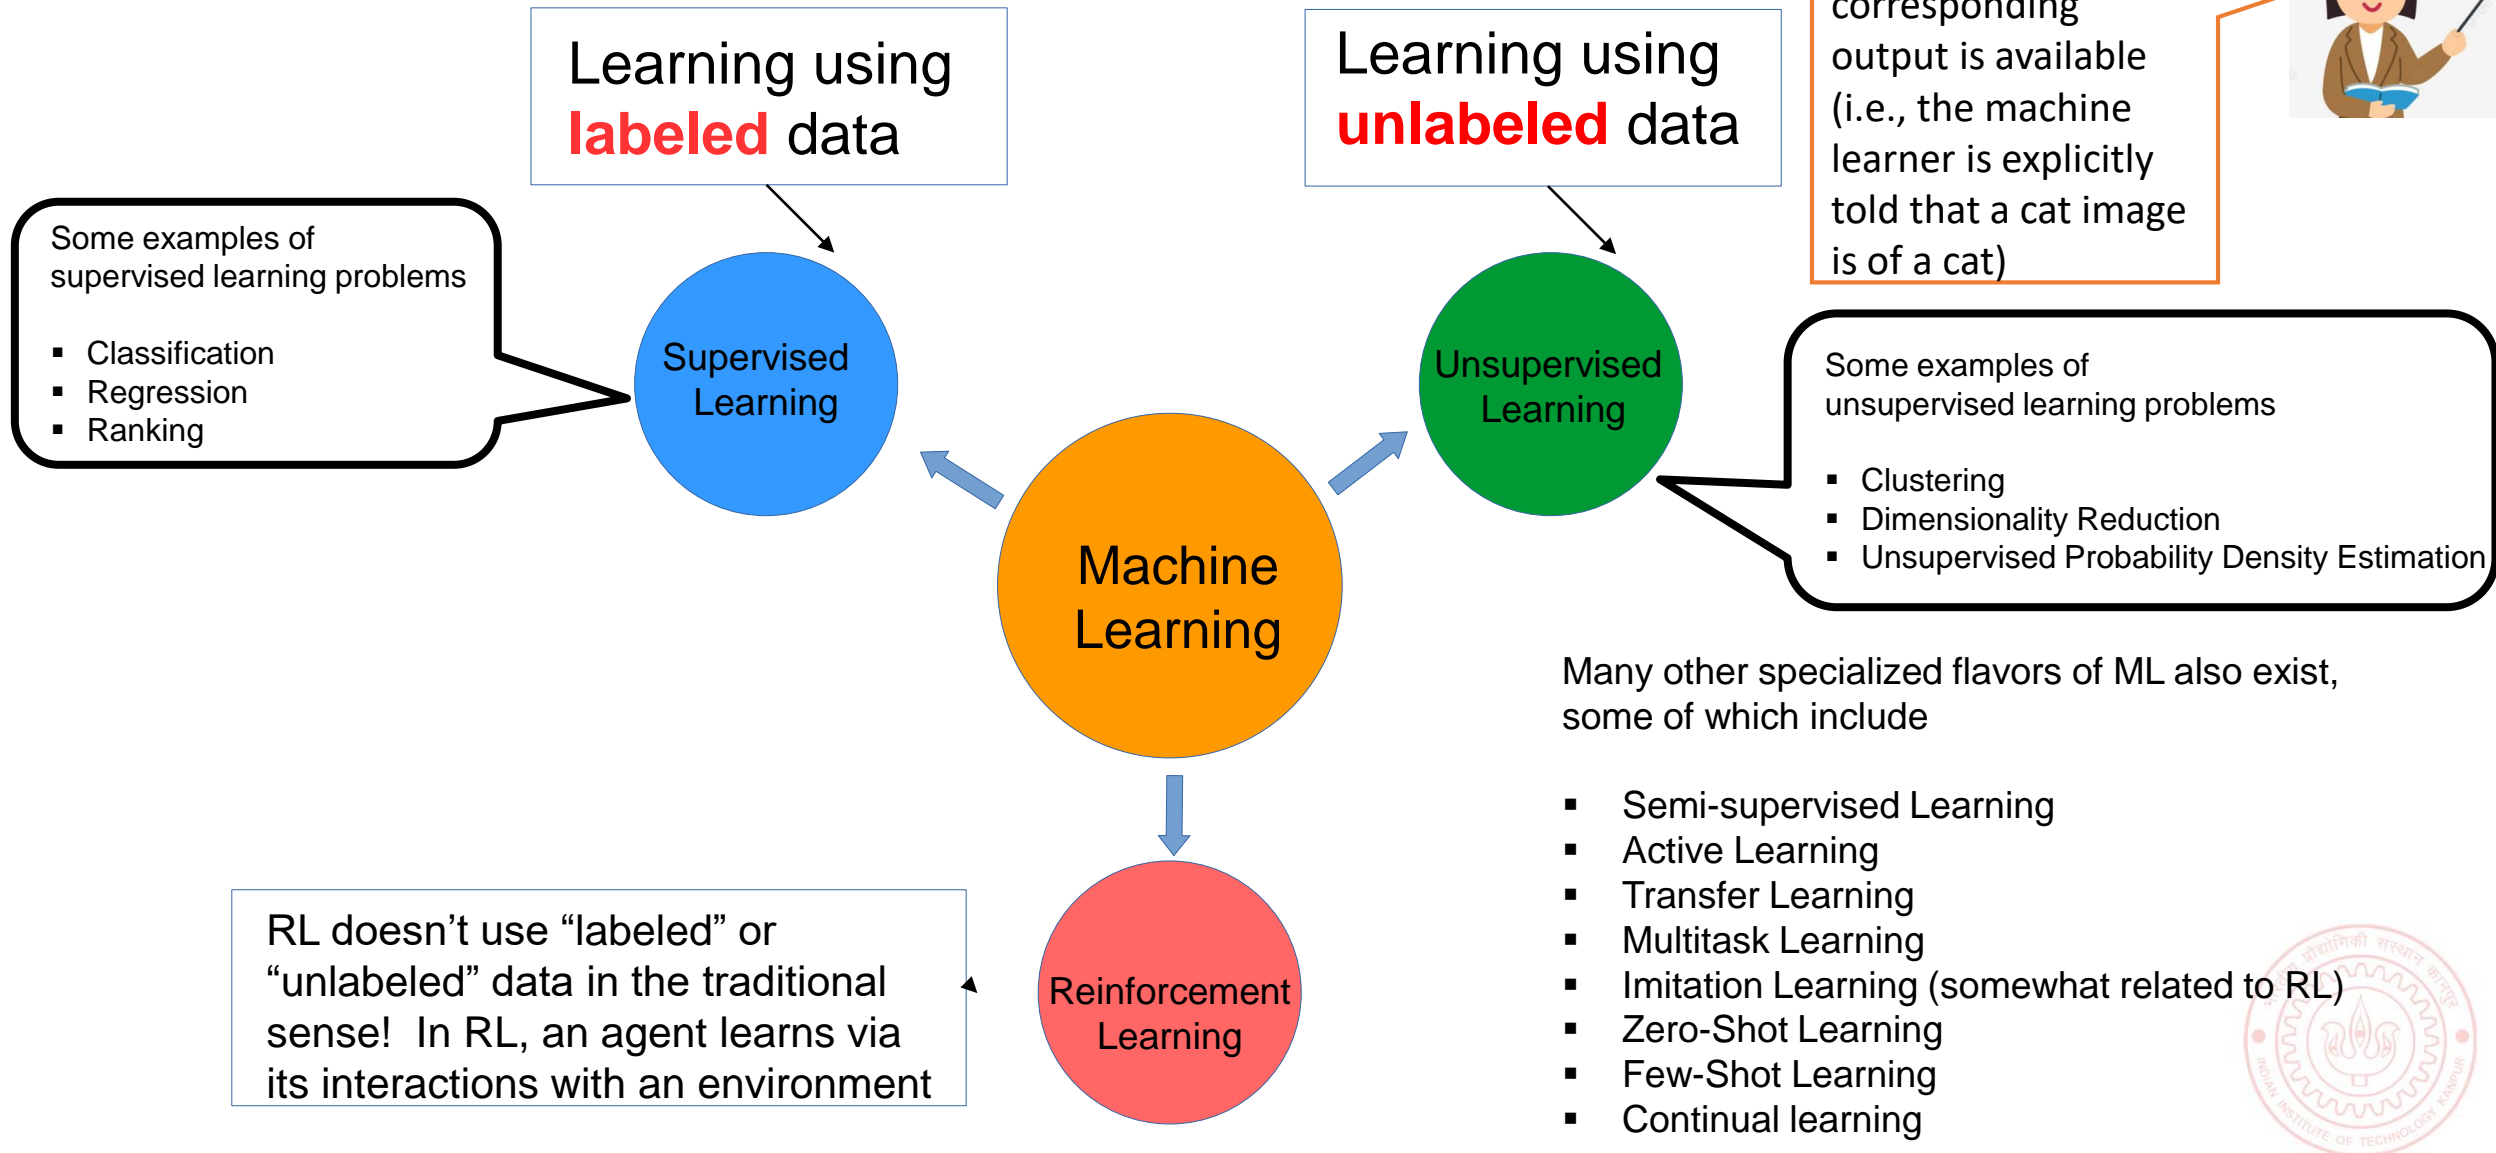
RL doesn't use “labeled” or “unlabeled” data in the traditional sense! In RL, an agent learns via its interactions with an environment

Reinforcement Learning



A Loose Taxonomy of ML

4



A Typical Supervised Learning Workflow




A Typical Supervised Learning Workflow

Note: This example is for the problem of **binary classification**, a supervised learning problem



A Typical Supervised Learning Workflow

5



Note: This example is for the problem of **binary classification**, a supervised learning problem



A Typical Supervised Learning Workflow

5

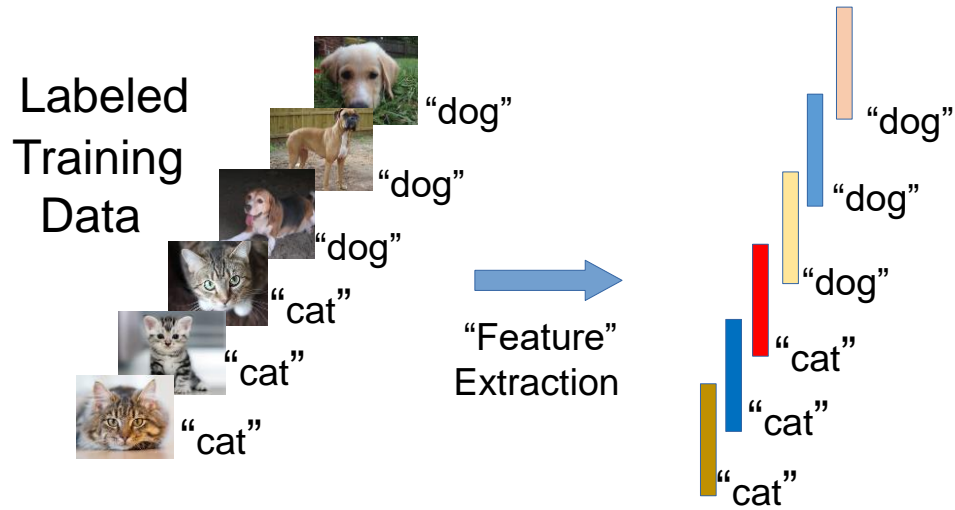


Note: This example is for the problem of **binary classification**, a supervised learning problem



A Typical Supervised Learning Workflow

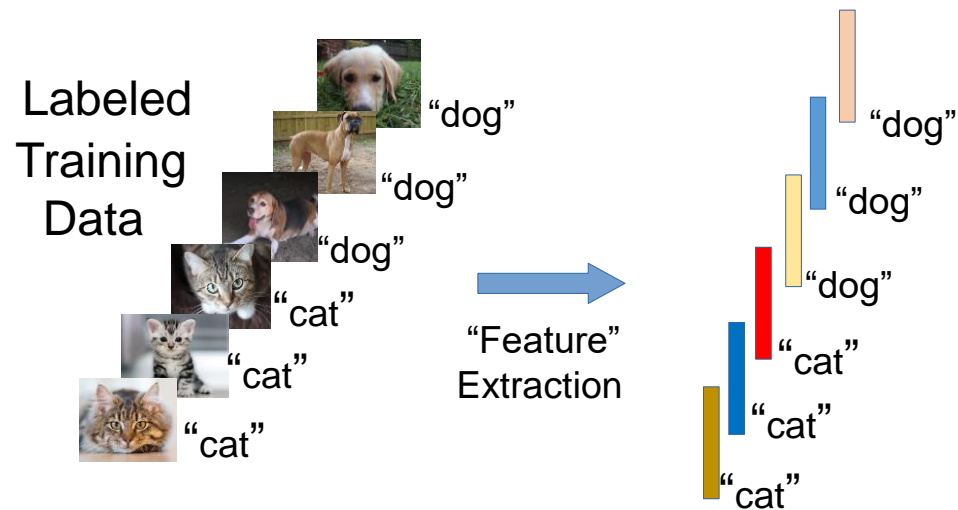
5




Note: This example is for the problem of **binary classification**, a supervised learning problem



A Typical Supervised Learning Workflow



Note: This example is for the problem of **binary classification**, a supervised learning problem


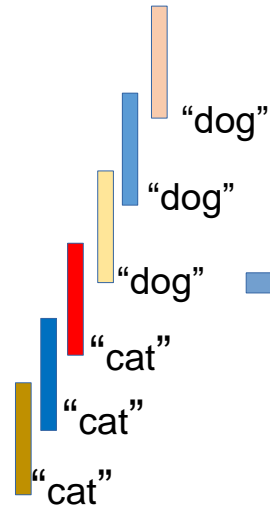
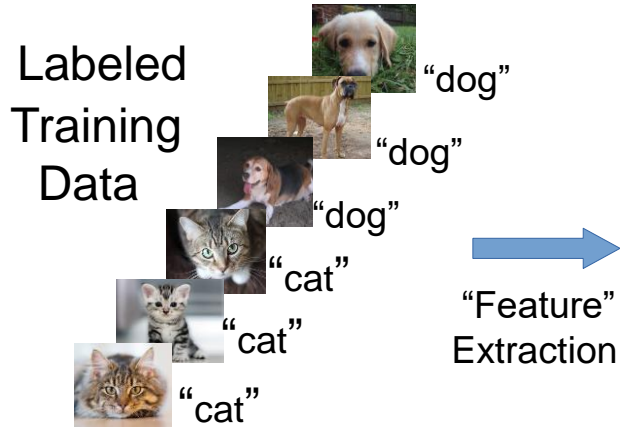


Feature extraction converts raw inputs to a **numeric representation** that the ML algo can understand and work with. More on feature extraction later.



A Typical Supervised Learning Workflow


5



Note: This example is for the problem of **binary classification**, a supervised learning problem


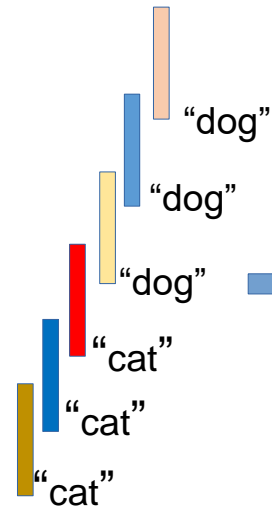
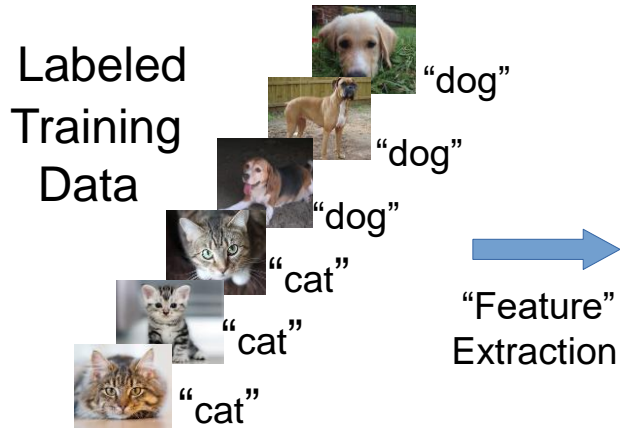


Feature extraction converts raw inputs to a **numeric representation** that the ML algo can understand and work with. More on feature extraction later.



A Typical Supervised Learning Workflow

5



Cat vs Dog
Prediction model

Note: This example is for the problem of **binary classification**, a supervised learning problem

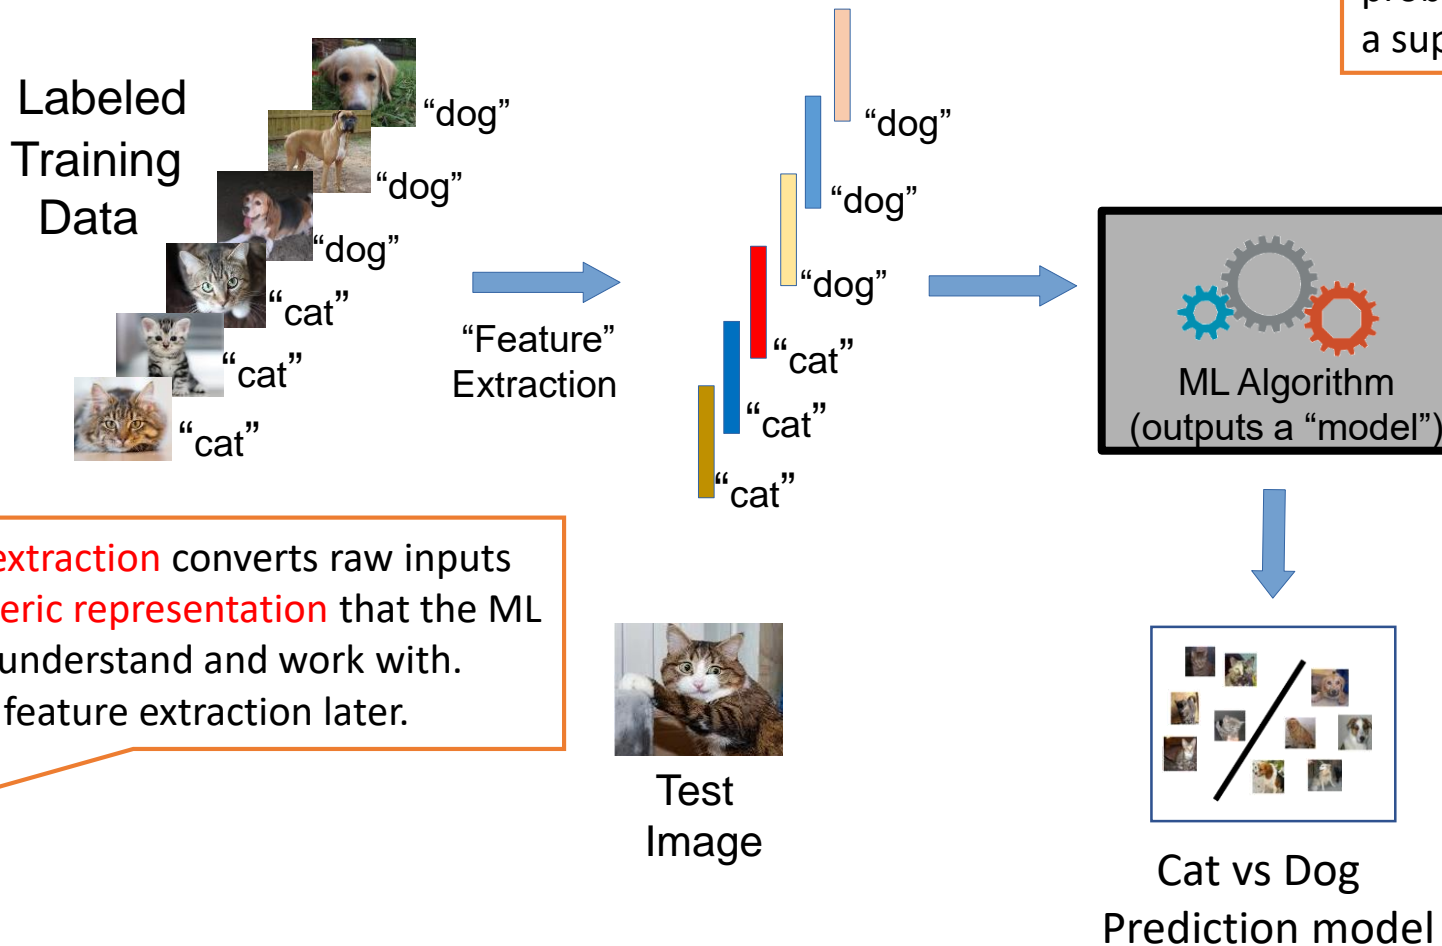


Feature extraction converts raw inputs to a **numeric representation** that the ML algo can understand and work with. More on feature extraction later.



A Typical Supervised Learning Workflow

5



Note: This example is for the problem of **binary classification**, a supervised learning problem

Feature extraction converts raw inputs to a **numeric representation** that the ML algo can understand and work with. More on feature extraction later.

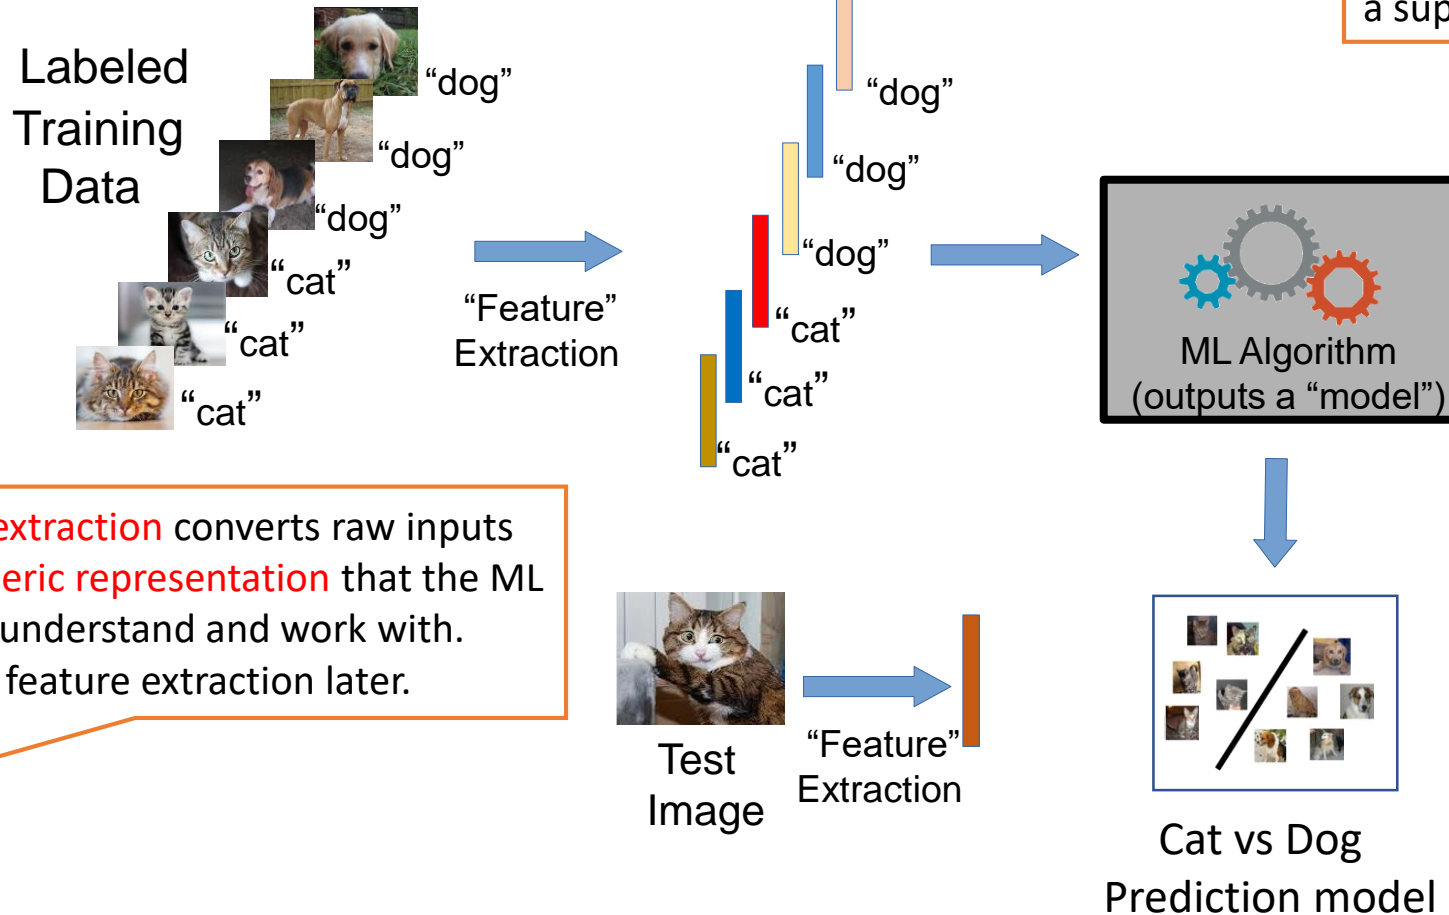


A Typical Supervised Learning Workflow


5



Note: This example is for the problem of **binary classification**, a supervised learning problem



Feature extraction converts raw inputs to a **numeric representation** that the ML algo can understand and work with. More on feature extraction later.

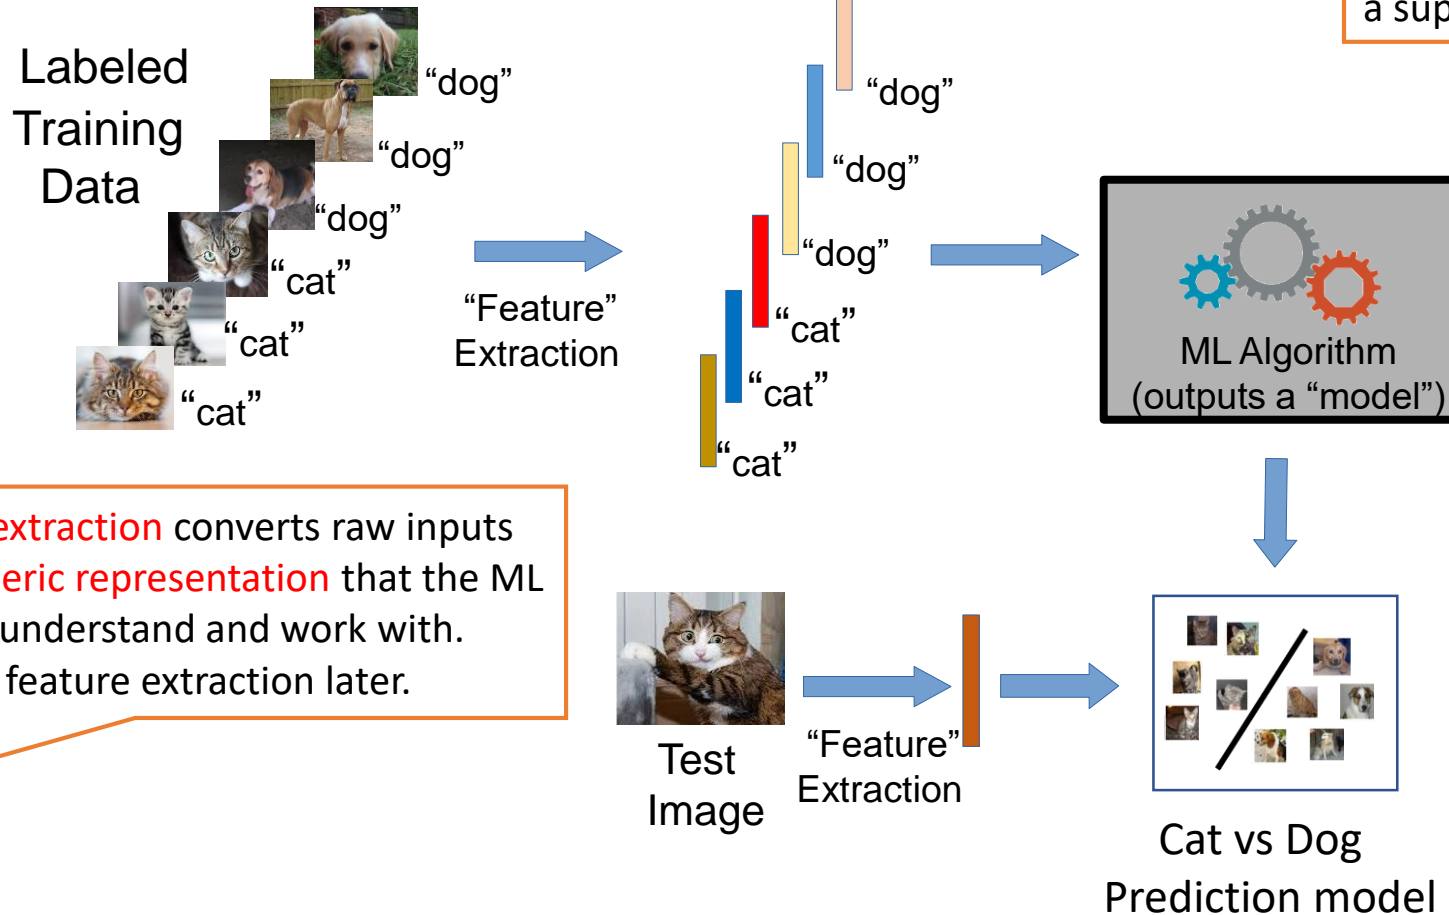


A Typical Supervised Learning Workflow

5



Note: This example is for the problem of **binary classification**, a supervised learning problem



Feature extraction converts raw inputs to a **numeric representation** that the ML algo can understand and work with. More on feature extraction later.




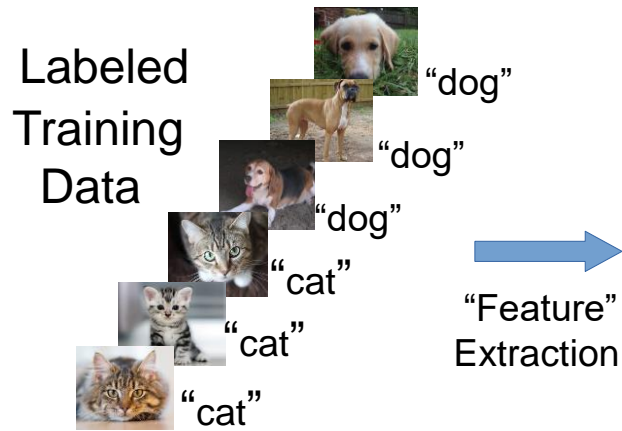


A Typical Supervised Learning Workflow

5



Note: This example is for the problem of **binary classification**, a supervised learning problem



Cat vs Dog
Prediction model

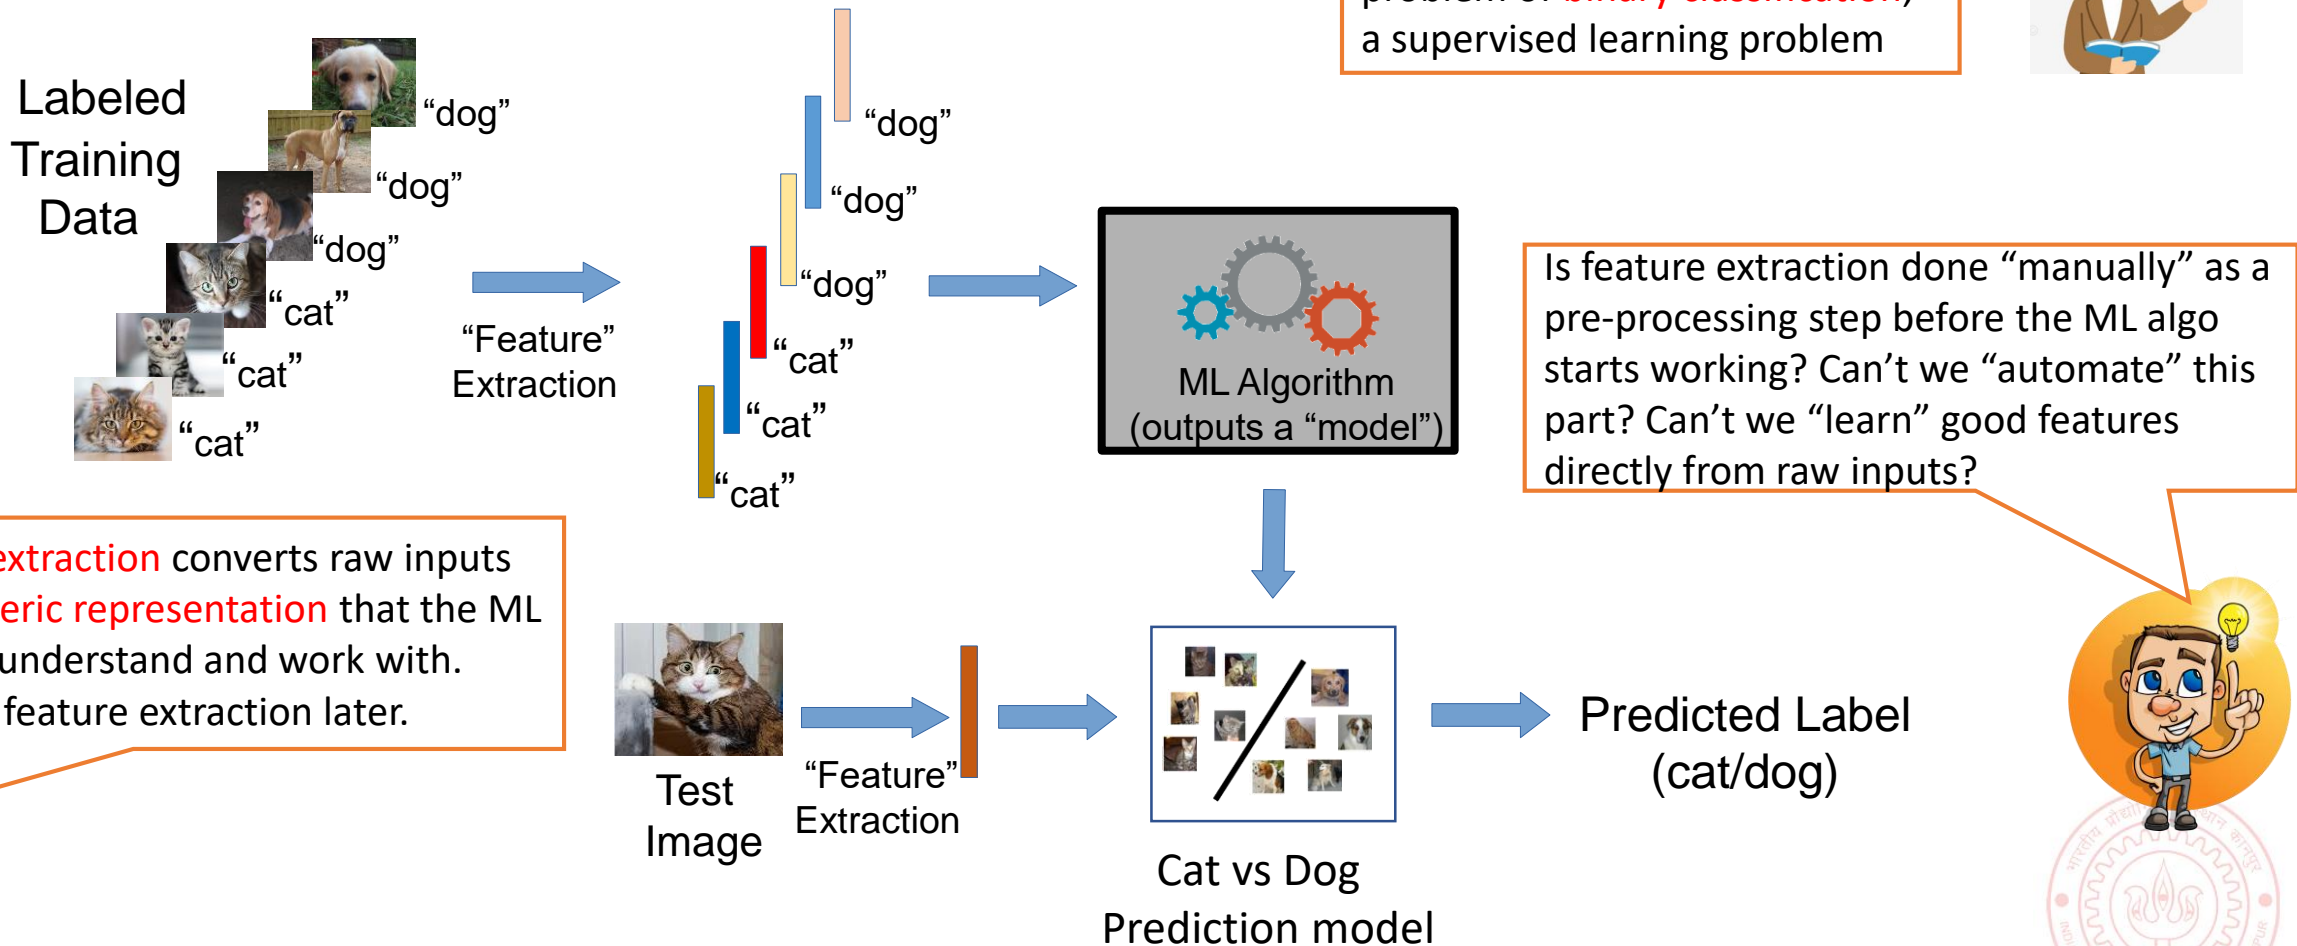
Predicted Label
(cat/dog)

Feature extraction converts raw inputs to a **numeric representation** that the ML algo can understand and work with. More on feature extraction later.




A Typical Supervised Learning Workflow



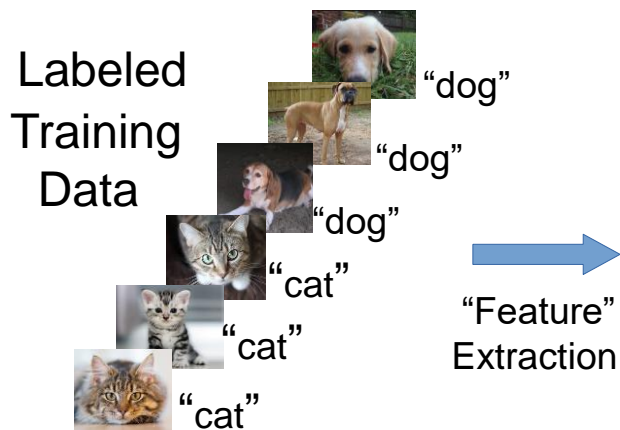
5



A Typical Supervised Learning Workflow



Note: This example is for the problem of **binary classification**, a supervised learning problem



Is feature extraction done “manually” as a pre-processing step before the ML algo starts working? Can’t we “automate” this part? Can’t we “learn” good features directly from raw inputs?

Feature extraction converts raw inputs to a **numeric representation** that the ML algo can understand and work with. More on feature extraction later.

Indeed. **Deep Learning** algos do precisely that! (**feature + model learning**). More on Deep Learning later.



Test Image




Cat vs Dog
Prediction model

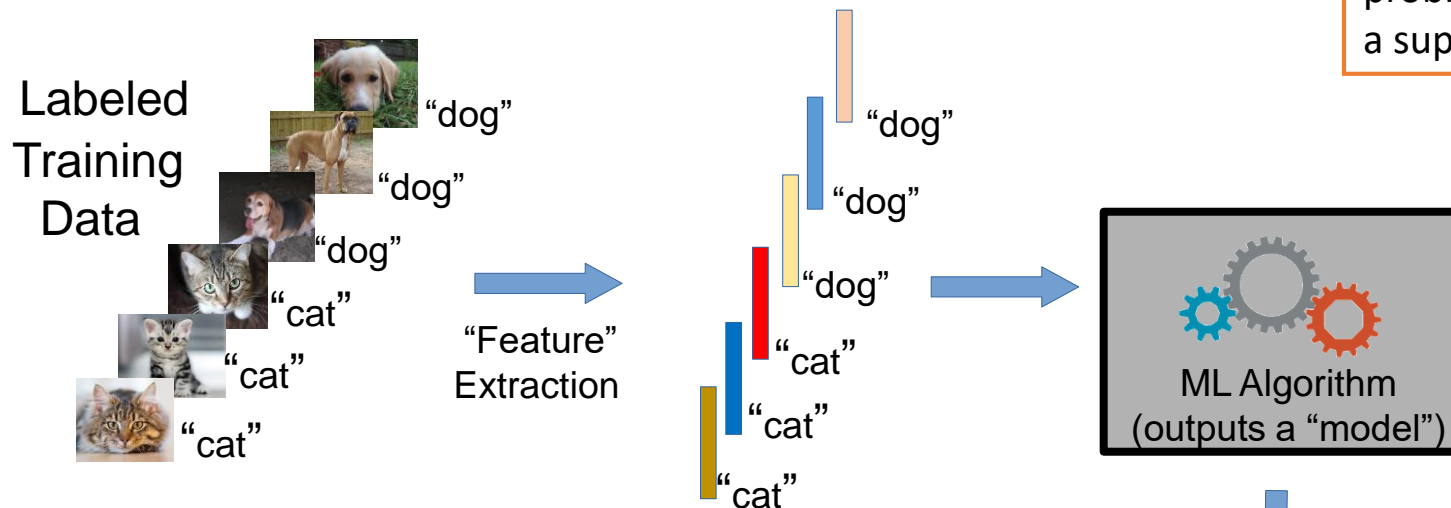
Predicted Label
(cat/dog)



A Typical Supervised Learning Workflow



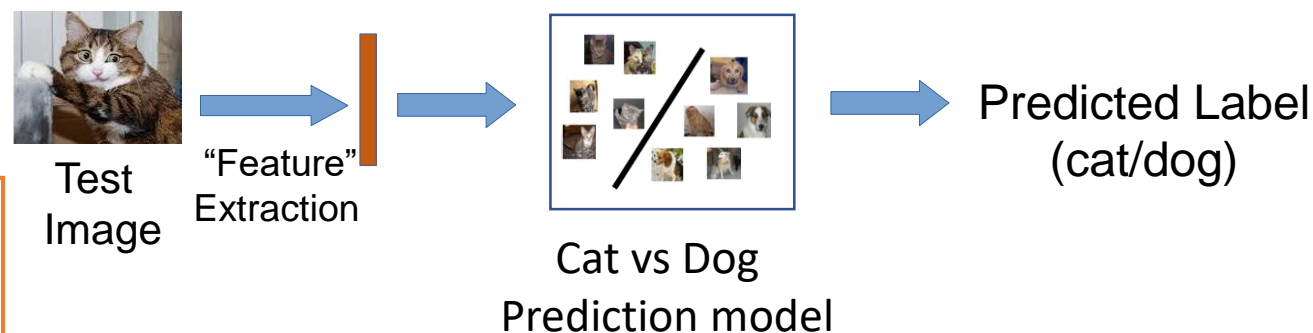
Note: This example is for the problem of **binary classification**, a supervised learning problem



Is feature extraction done “manually” as a pre-processing step before the ML algo starts working? Can’t we “automate” this part? Can’t we “learn” good features directly from raw inputs?

Feature extraction converts raw inputs to a **numeric representation** that the ML algo can understand and work with. More on feature extraction later.

Indeed. **Deep Learning** algos do precisely that! (**feature + model learning**). More on Deep Learning later.



A Typical Unsupervised Learning Workflow



A Typical Unsupervised Learning Workflow

Note: This example is for the problem of **data clustering**, an unsupervised learning problem



A Typical Unsupervised Learning Workflow

Unlabeled
Data

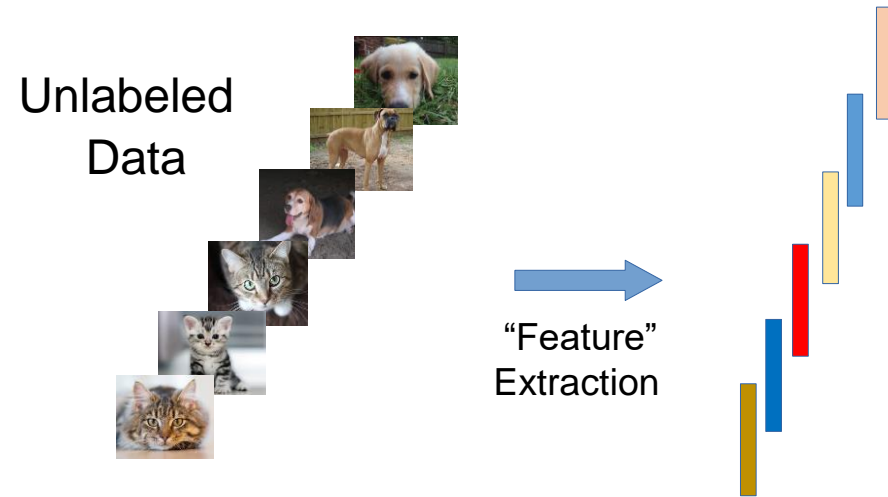


Note: This example is for the problem of **data clustering**, an unsupervised learning problem



A Typical Unsupervised Learning Workflow

6




Note: This example is for the problem of **data clustering**, an unsupervised learning problem



A Typical Unsupervised Learning Workflow

6

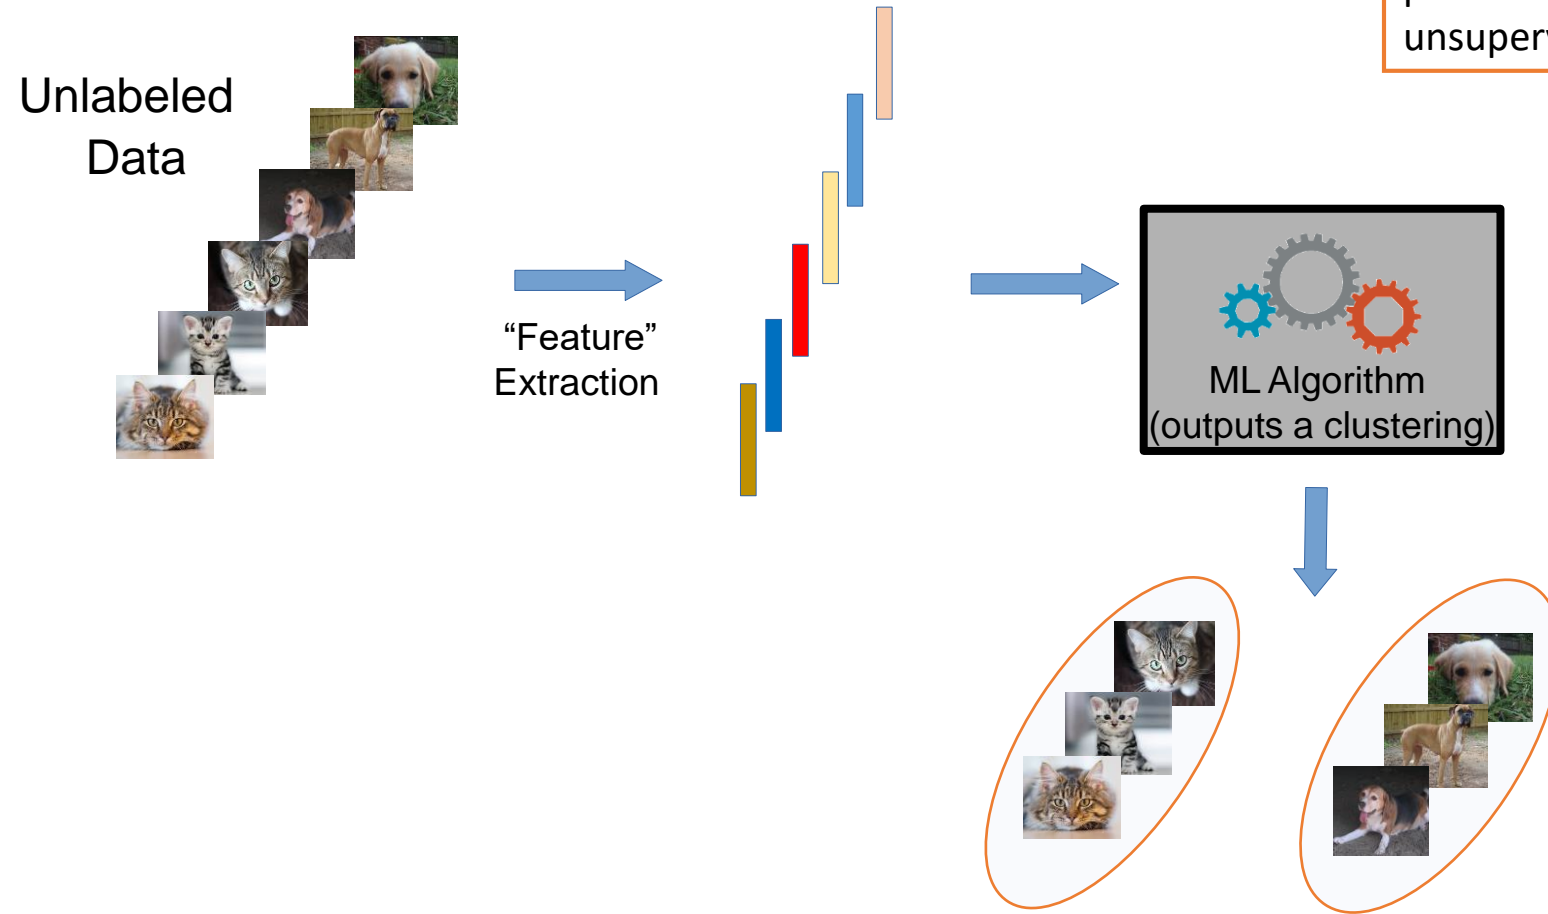


Note: This example is for the problem of **data clustering**, an unsupervised learning problem



A Typical Unsupervised Learning Workflow

6

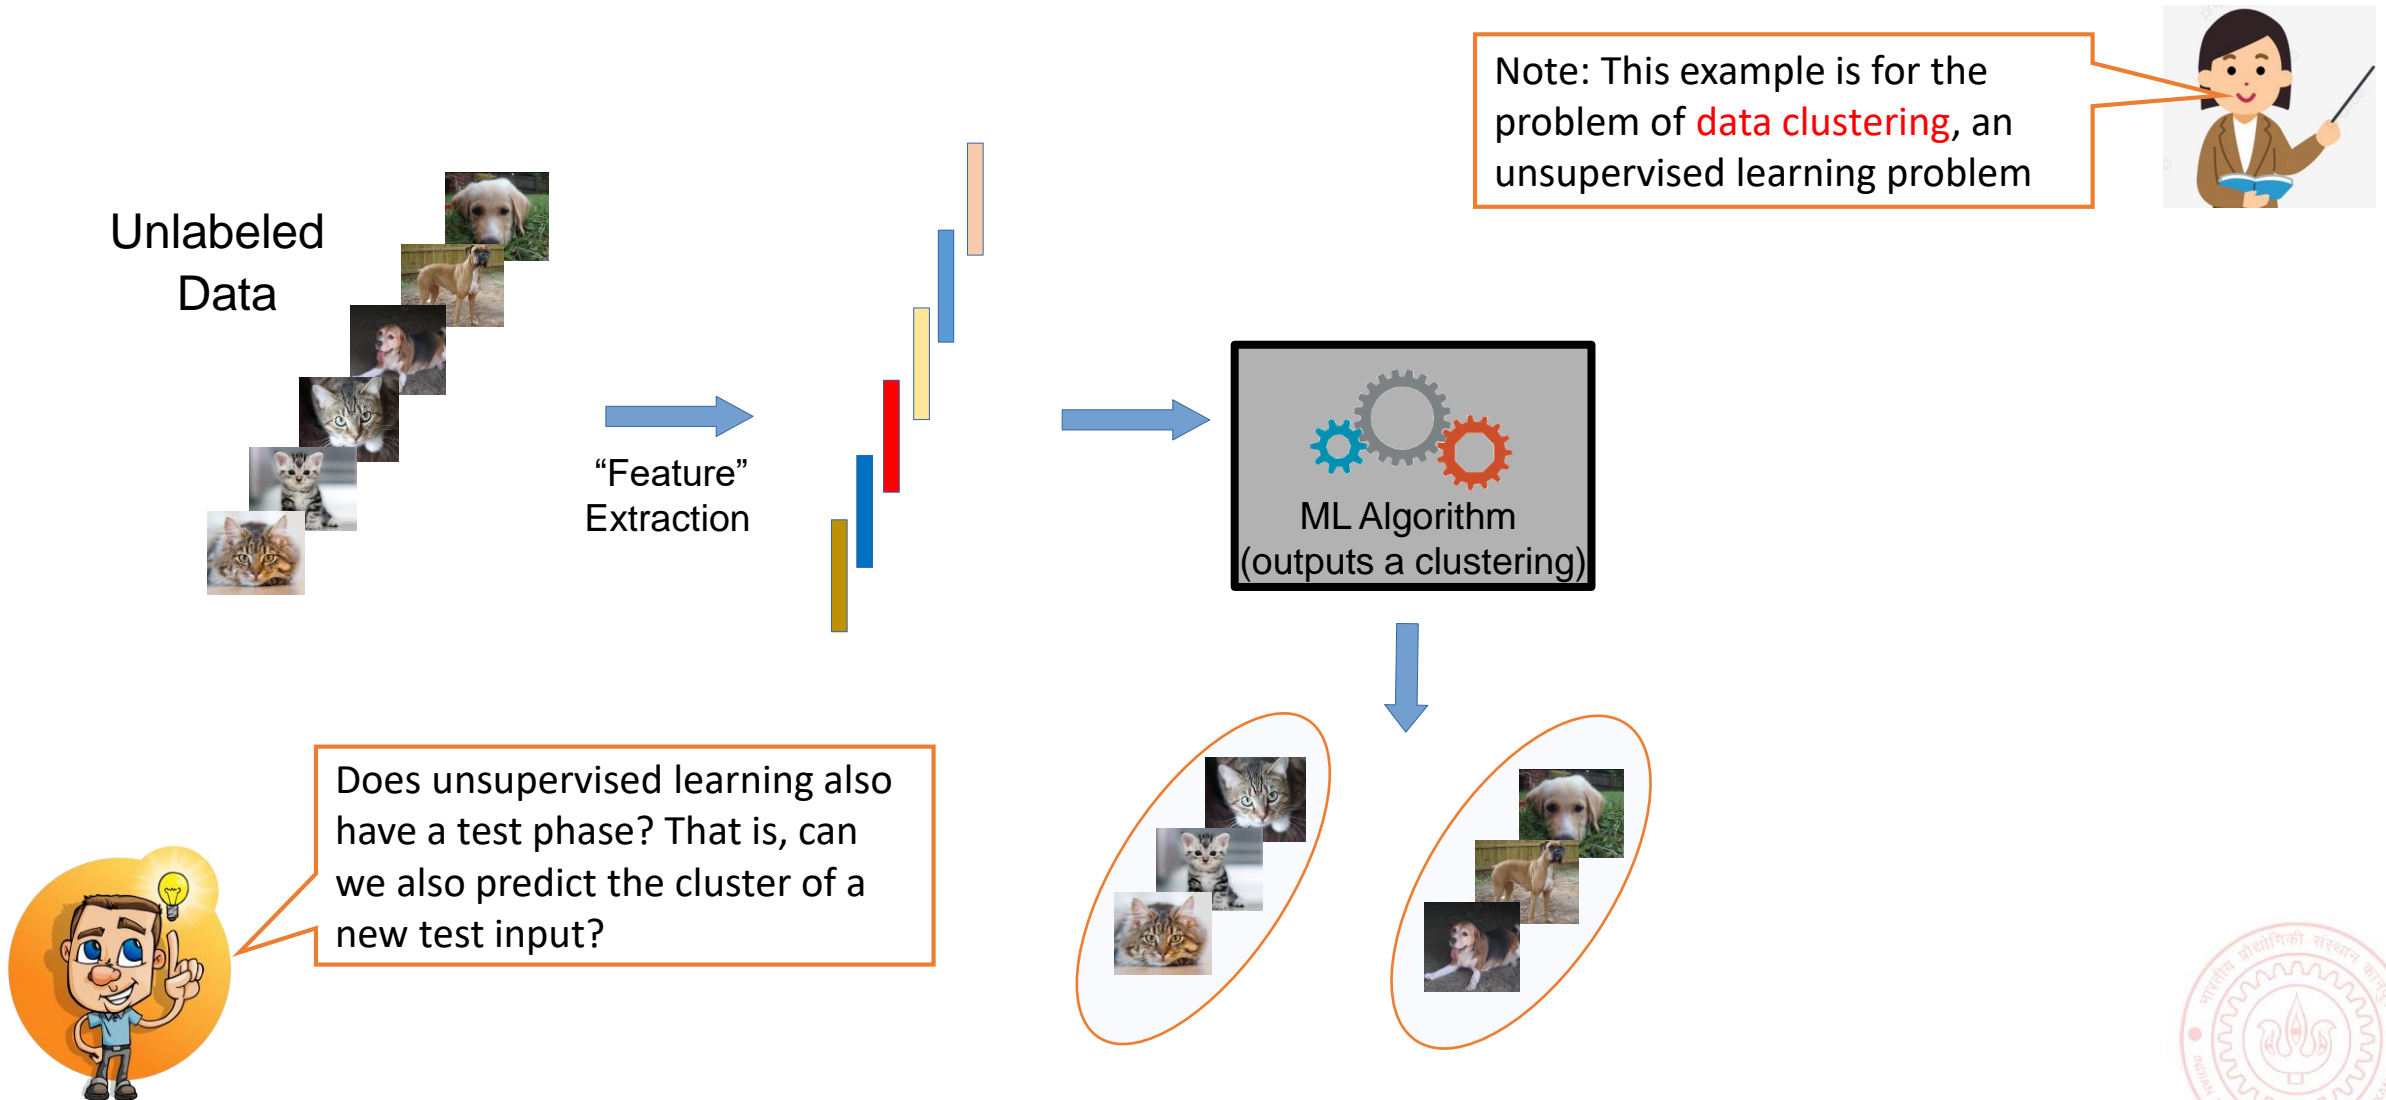


Note: This example is for the problem of **data clustering**, an unsupervised learning problem



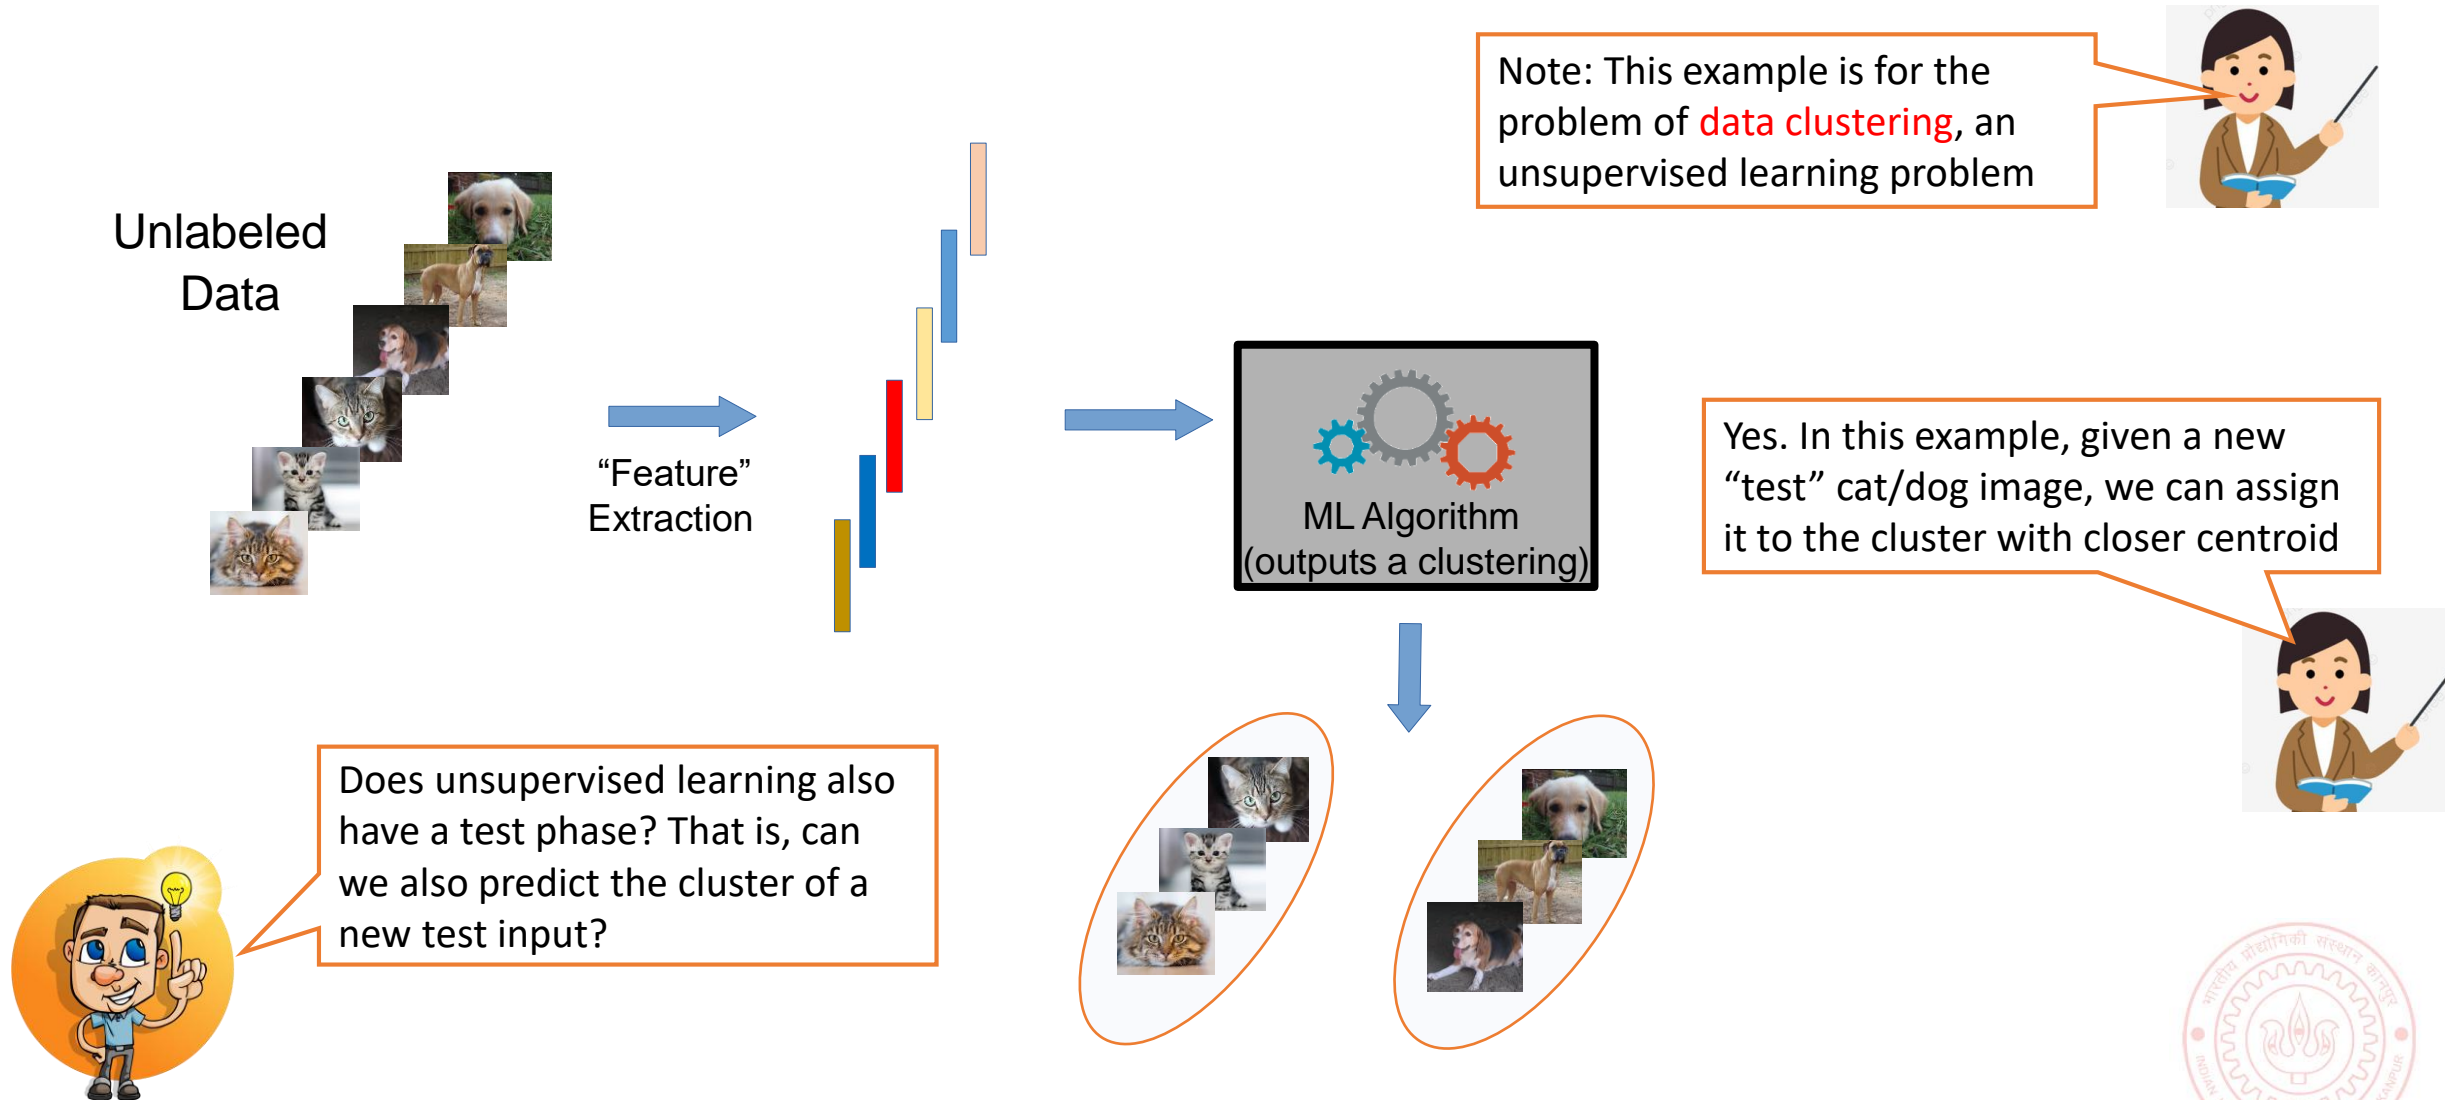
A Typical Unsupervised Learning Workflow

6



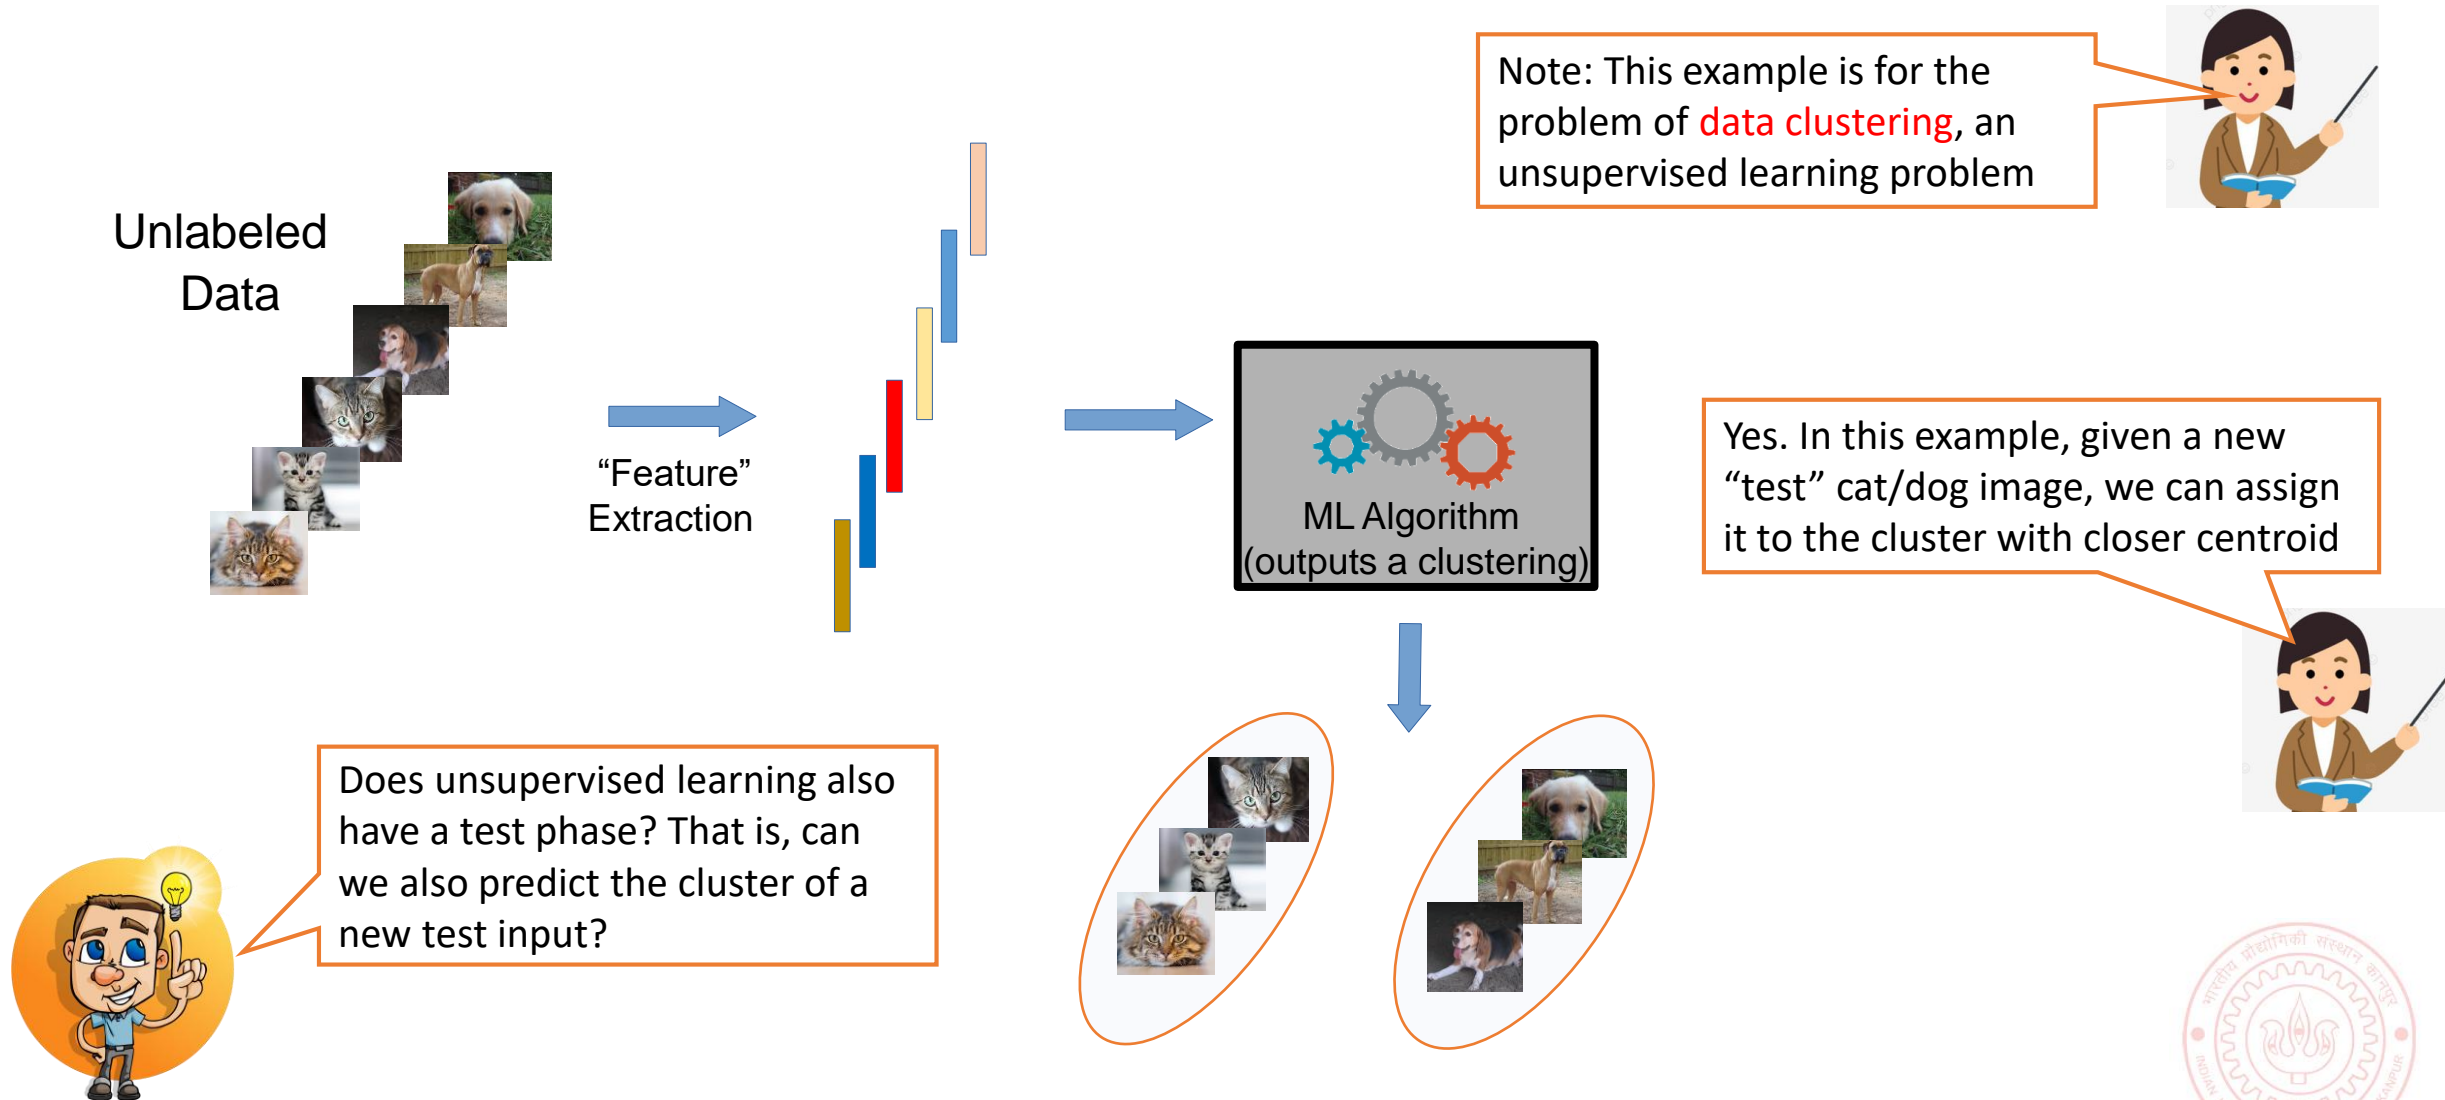
A Typical Unsupervised Learning Workflow

6



A Typical Unsupervised Learning Workflow

6



A Typical Reinforcement Learning Workflow

7



A Typical Reinforcement Learning Workflow

7

Wish to teach an agent optimal policy for some task

Agent does the following repeatedly


- Senses/observes the environment
- Takes an action based on its current policy
- Receives a reward for that action
- Updates its policy

Agent's goal is to maximize its overall reward



A Typical Reinforcement Learning Workflow

7



Wish to teach an agent optimal policy for some task

Agent does the following repeatedly

- Senses/observes the environment
- Takes an action based on its current policy
- Receives a reward for that action
- Updates its policy

Agent's goal is to maximize its overall reward



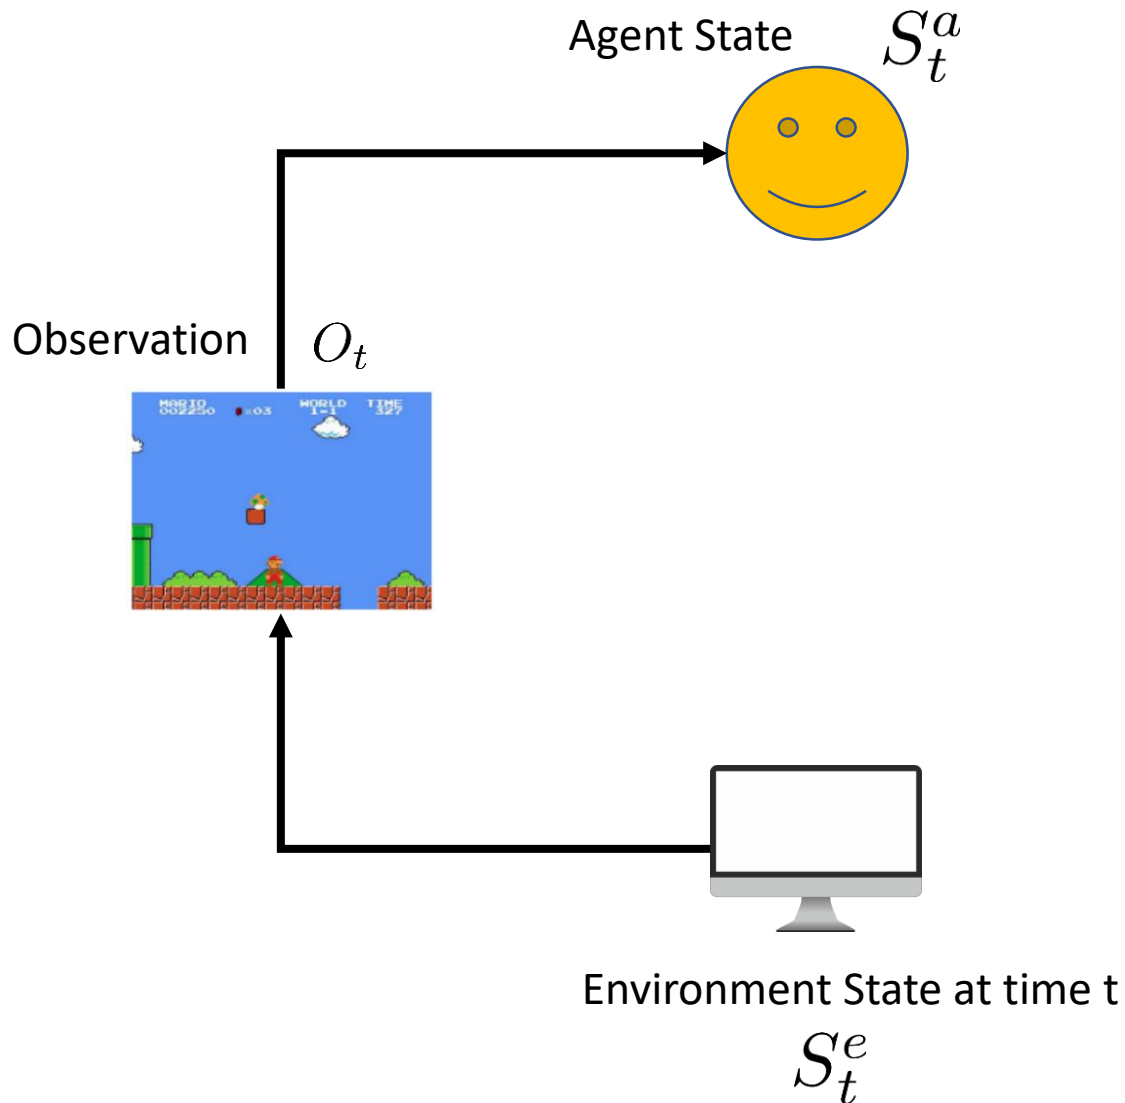
Environment State at time t

S_t^e



A Typical Reinforcement Learning Workflow

7



Wish to teach an agent optimal policy for some task

Agent does the following repeatedly

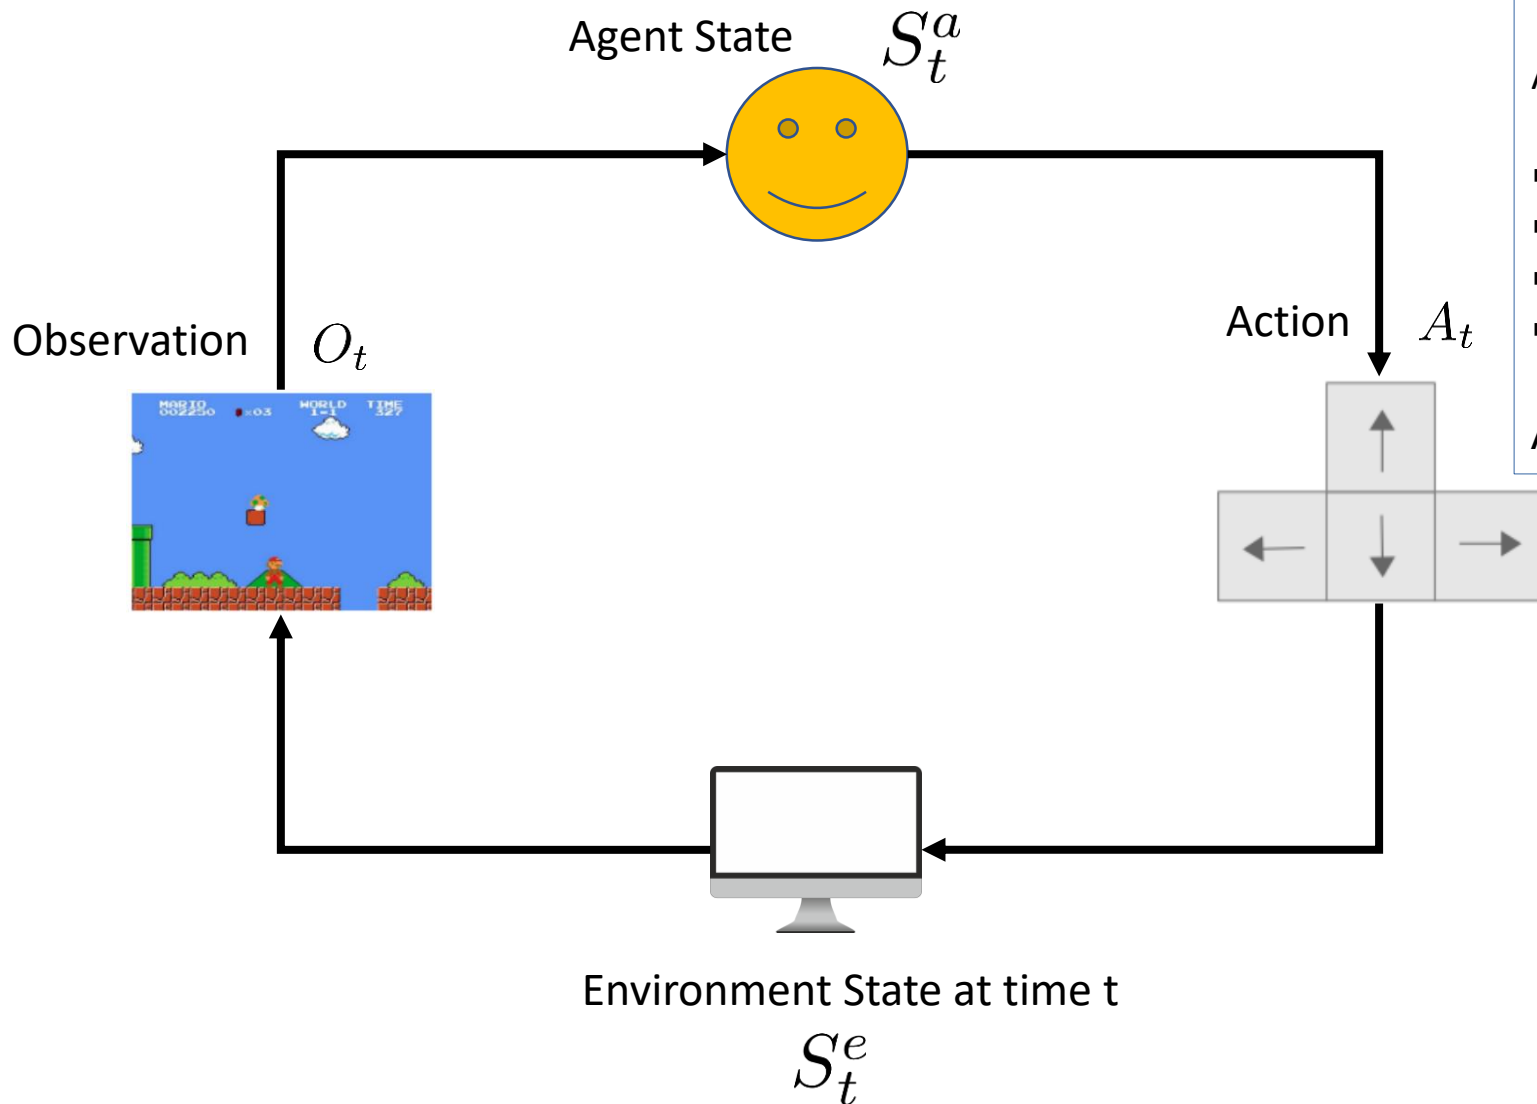
- Senses/observes the environment
- Takes an action based on its current policy
- Receives a reward for that action
- Updates its policy

Agent's goal is to maximize its overall reward



A Typical Reinforcement Learning Workflow

7



Wish to teach an agent optimal policy for some task

Agent does the following repeatedly

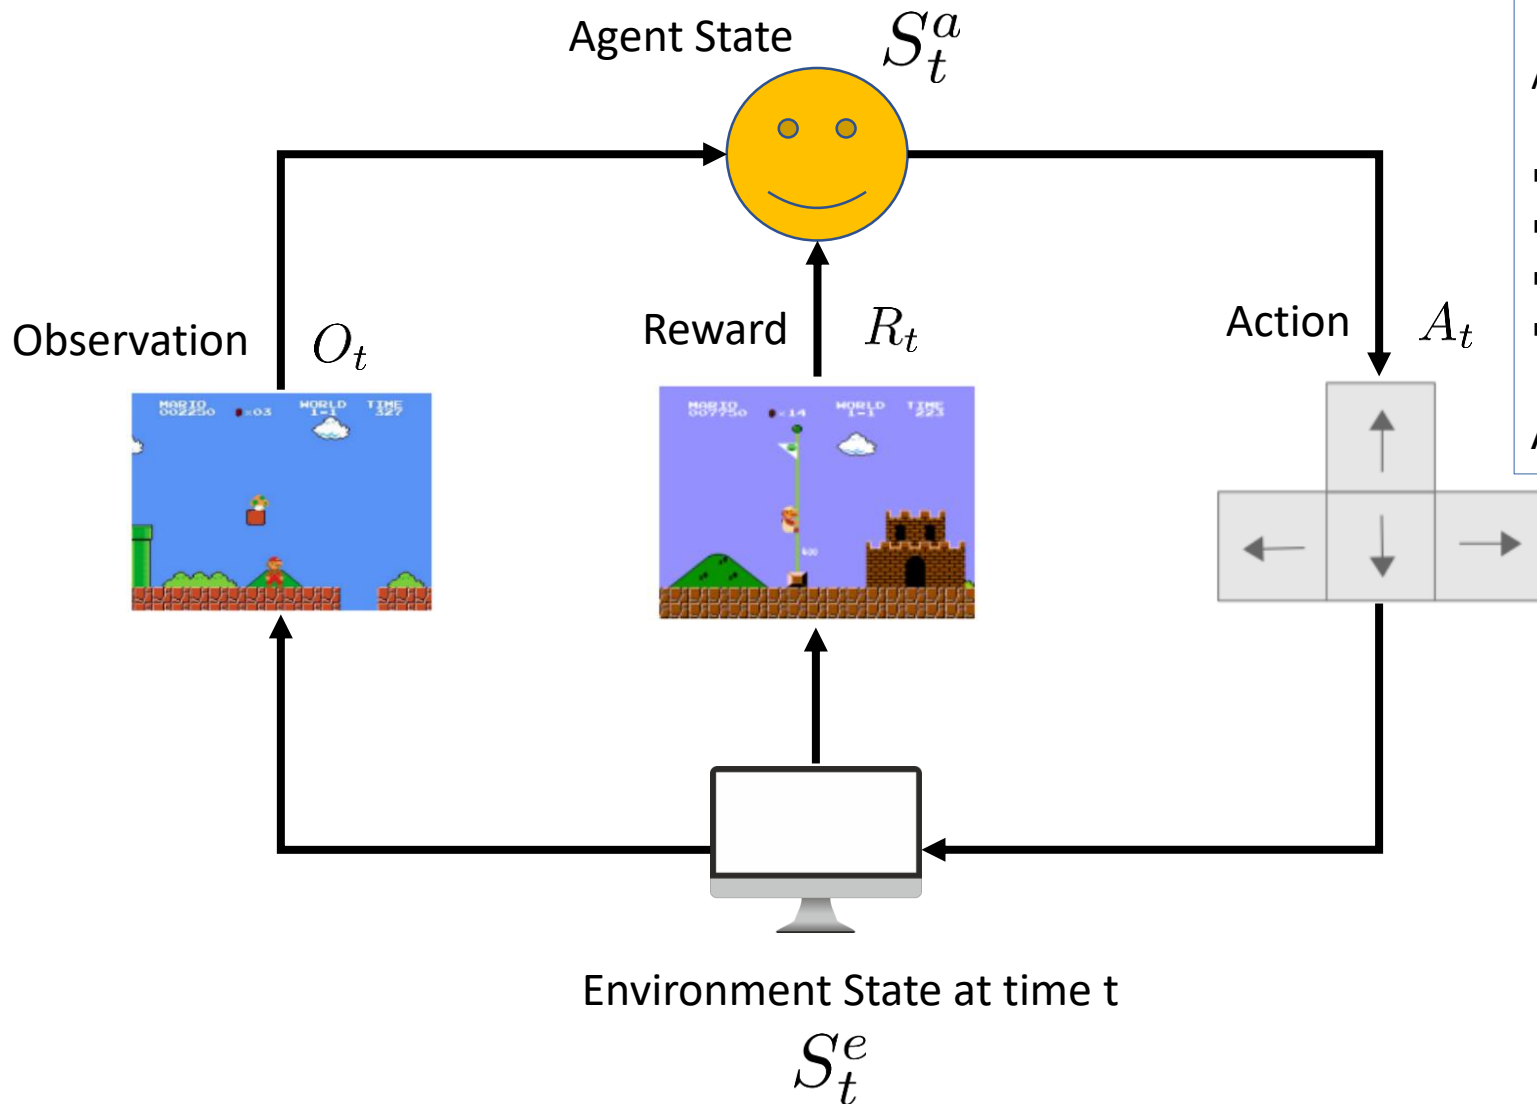
- Senses/observes the environment
- Takes an action based on its current policy
- Receives a reward for that action
- Updates its policy

Agent's goal is to maximize its overall reward



A Typical Reinforcement Learning Workflow

7



Wish to teach an agent optimal policy for some task

Agent does the following repeatedly


- Senses/observes the environment
- Takes an action based on its current policy
- Receives a reward for that action
- Updates its policy

Agent's goal is to maximize its overall reward



A Typical Reinforcement Learning Workflow

7



Wish to teach an agent optimal policy for some task

Agent does the following repeatedly

- Senses/observes the environment
- Takes an action based on its current policy
- Receives a reward for that action
- Updates its policy

Agent's goal is to maximize its overall reward

There IS supervision, not explicit
(as in Supervised Learning) but
rather implicit (feedback based)



ML: Some Perspectives



Geometric Perspective



Geometric Perspective

- Basic fact: Inputs in ML problems can often be represented as **points or vectors** in some vector space



Geometric Perspective

9

Recall that feature extraction converts inputs into a **numeric representation**

- Basic fact: Inputs in ML problems can often be represented as **points or vectors** in some vector space



Geometric Perspective

9

Recall that feature extraction converts inputs into a **numeric representation**

- Basic fact: Inputs in ML problems can often be represented as **points or vectors** in some vector space
- Doing ML on such data can thus be seen from a geometric view



Geometric Perspective

9

Recall that feature extraction converts inputs into a **numeric representation**

- Basic fact: Inputs in ML problems can often be represented as **points or vectors** in some vector space
- Doing ML on such data can thus be seen from a geometric view

Regression: A supervised learning problem. Goal is to model the relationship between input (x) and real-valued output (y). This is akin to a **line or curve fitting** problem




Geometric Perspective

9

Recall that feature extraction converts inputs into a **numeric representation**

- Basic fact: Inputs in ML problems can often be represented as **points or vectors** in some vector space
- Doing ML on such data can thus be seen from a geometric view

Regression: A supervised learning problem. Goal is to model the relationship between input (x) and real-valued output (y). This is akin to a **line or curve fitting** problem





Geometric Perspective

9

Recall that feature extraction converts inputs into a **numeric representation**

- Basic fact: Inputs in ML problems can often be represented as **points or vectors** in some vector space
- Doing ML on such data can thus be seen from a geometric view

Regression: A supervised learning problem. Goal is to model the relationship between input (x) and real-valued output (y). This is akin to a **line or curve fitting** problem





Geometric Perspective

9

Recall that feature extraction converts inputs into a **numeric representation**

- Basic fact: Inputs in ML problems can often be represented as **points or vectors** in some vector space
- Doing ML on such data can thus be seen from a geometric view

Regression: A supervised learning problem. Goal is to model the relationship between input (x) and real-valued output (y). This is akin to a **line or curve fitting** problem



Classification: A supervised learning problem. Goal is to learn a to predict which of the two or more classes an input belongs to. Akin to learning **linear/nonlinear separator** for the inputs





Geometric Perspective

9





Recall that feature extraction converts inputs into a **numeric representation**

- Basic fact: Inputs in ML problems can often be represented as **points or vectors** in some vector space
- Doing ML on such data can thus be seen from a geometric view

Regression: A supervised learning problem. Goal is to model the relationship between input (x) and real-valued output (y). This is akin to a **line or curve fitting** problem



Classification: A supervised learning problem. Goal is to learn a to predict which of the two or more classes an input belongs to. Akin to learning **linear/nonlinear separator** for the inputs





Geometric Perspective

9





Recall that feature extraction converts inputs into a **numeric representation**

- Basic fact: Inputs in ML problems can often be represented as **points or vectors** in some vector space
- Doing ML on such data can thus be seen from a geometric view

Regression: A supervised learning problem. Goal is to model the relationship between input (x) and real-valued output (y). This is akin to a **line or curve fitting** problem



Classification: A supervised learning problem. Goal is to learn a to predict which of the two or more classes an input belongs to. Akin to learning **linear/nonlinear separator** for the inputs



Geometric Perspective




Geometric Perspective

Clustering: An unsupervised learning problem. Goal is to group inputs in a few clusters **based on their similarities with each other**




Geometric Perspective

Clustering: An unsupervised learning problem. Goal is to group inputs in a few clusters **based on their similarities with each other**



Geometric Perspective

Clustering: An unsupervised learning problem. Goal is to group inputs in a few clusters **based on their similarities with each other**




Clustering looks like classification to me. Is there any difference?



Geometric Perspective

Clustering: An unsupervised learning problem. Goal is to group inputs in a few clusters **based on their similarities with each other**



Clustering looks like classification to me. Is there any difference?




Yes. In clustering, we don't know the labels. Goal is to separate them without any labeled "supervision"



Geometric Perspective

Clustering: An unsupervised learning problem. Goal is to group inputs in a few clusters **based on their similarities with each other**



Clustering looks like classification to me. Is there any difference?



Yes. In clustering, we don't know the labels. Goal is to separate them without any labeled "supervision"




Dimensionality Reduction: An unsupervised learning problem. Goal is to **compress the size** of each input without losing much information present in the data



Geometric Perspective

Clustering: An unsupervised learning problem. Goal is to group inputs in a few clusters **based on their similarities with each other**




Clustering looks like classification to me. Is there any difference?



Yes. In clustering, we don't know the labels. Goal is to separate them without any labeled "supervision"




Dimensionality Reduction: An unsupervised learning problem. Goal is to **compress the size** of each input without losing much information present in the data



Geometric Perspective

10

Clustering: An unsupervised learning problem. Goal is to group inputs in a few clusters **based on their similarities with each other**





Clustering looks like classification to me. Is there any difference?



Yes. In clustering, we don't know the labels. Goal is to separate them without any labeled "supervision"



Dimensionality Reduction: An unsupervised learning problem. Goal is to **compress the size** of each input without losing much information present in the data



Three-dim to two-dim
nonlinear projection
(a.k.a. manifold learning)



Perspective as function approximation



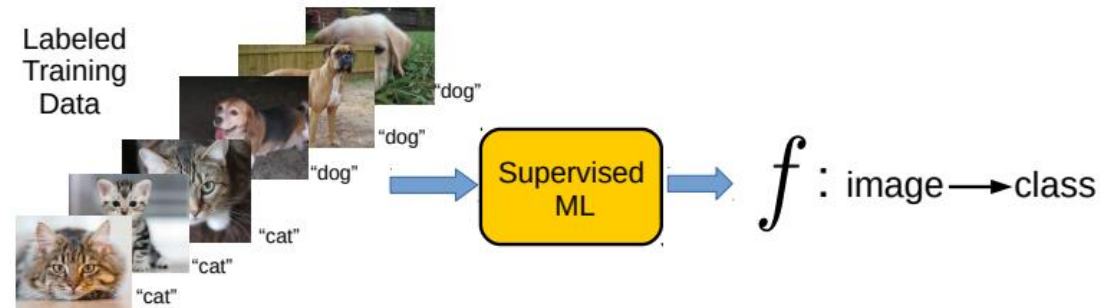
Perspective as function approximation

- Supervised Learning (“predict output given input”) can be usually thought of as learning a **function** f that maps each input to the corresponding output




Perspective as function approximation

- Supervised Learning (“predict output given input”) can be usually thought of as learning a **function** f that maps each input to the corresponding output



Perspective as function approximation

- Supervised Learning (“predict output given input”) can be usually thought of as learning a **function** f that maps each input to the corresponding output




- Unsupervised Learning (“model/compress inputs”) can also be usually thought of as learning a **function** f that maps each input to a compact representation



Perspective as function approximation

- Supervised Learning (“predict output given input”) can be usually thought of as learning a **function** f that maps each input to the corresponding output




- Unsupervised Learning (“model/compress inputs”) can also be usually thought of as learning a **function** f that maps each input to a compact representation

Harder since we don't know the labels in this case



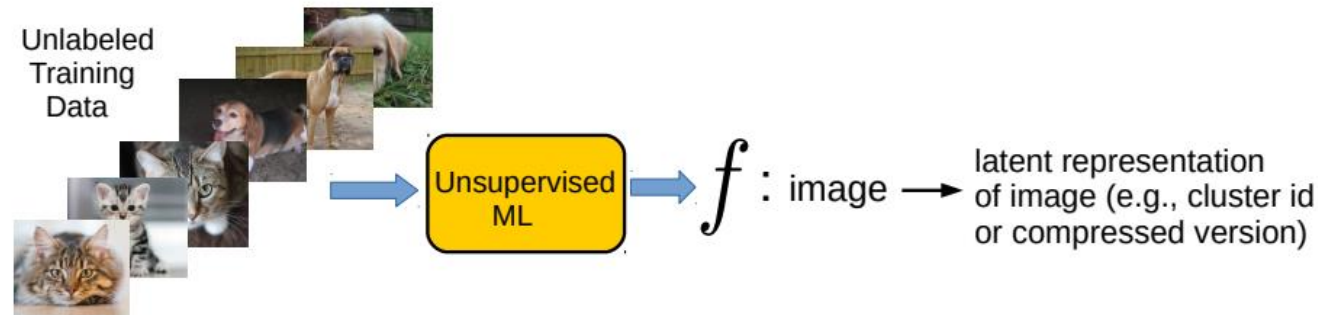
Perspective as function approximation

- Supervised Learning (“predict output given input”) can be usually thought of as learning a **function** f that maps each input to the corresponding output



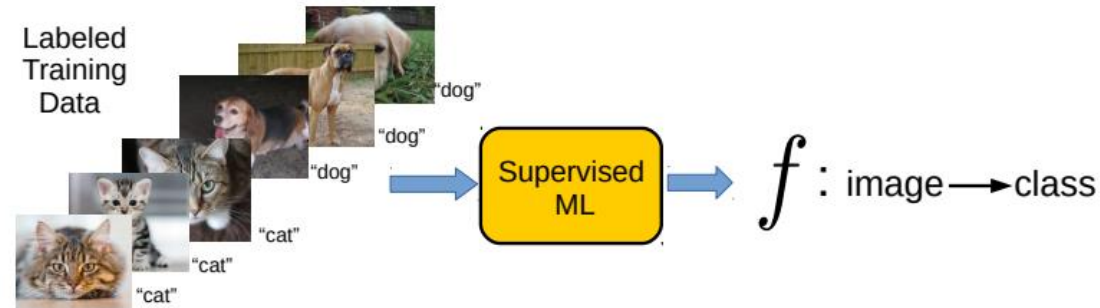
- Unsupervised Learning (“model/compress inputs”) can also be usually thought of as learning a **function** f that maps each input to a compact representation

Harder since we don't know the labels in this case




Perspective as function approximation

- Supervised Learning (“predict output given input”) can be usually thought of as learning a **function** f that maps each input to the corresponding output



- Unsupervised Learning (“model/compress inputs”) can also be usually thought of as learning a **function** f that maps each input to a compact representation

Harder since we don't know the labels in this case



- Reinforcement Learning can also be seen as doing function approximation



Perspective as probability estimation



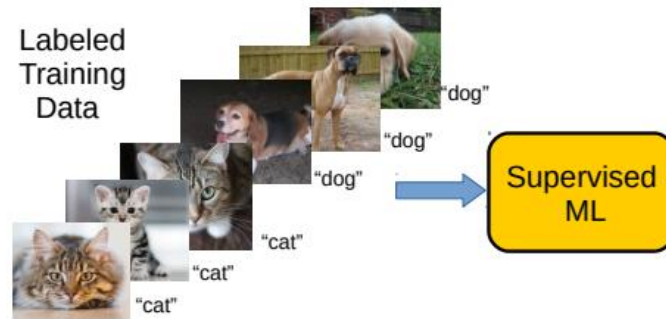
Perspective as probability estimation

- Supervised Learning (“predict output given input”) can be thought of as estimating the **conditional probability** of each possible output given an input



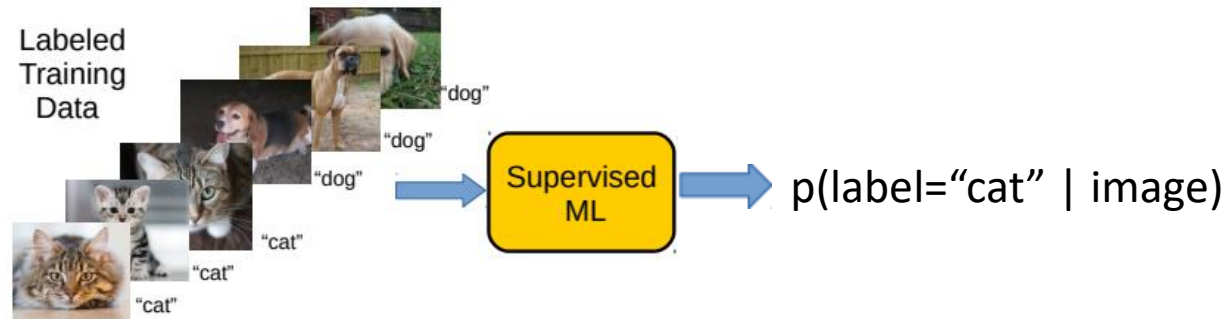
Perspective as probability estimation

- Supervised Learning (“predict output given input”) can be thought of as estimating the **conditional probability** of each possible output given an input



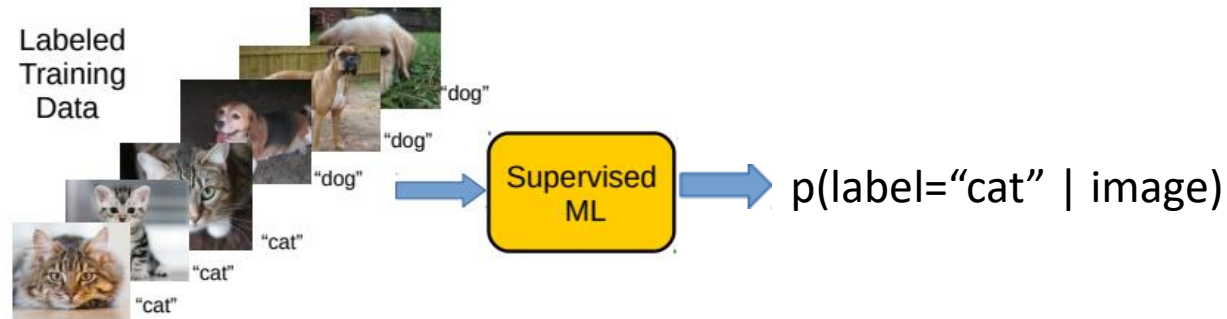
Perspective as probability estimation

- Supervised Learning (“predict output given input”) can be thought of as estimating the **conditional probability** of each possible output given an input



Perspective as probability estimation

- Supervised Learning (“predict output given input”) can be thought of as estimating the **conditional probability** of each possible output given an input

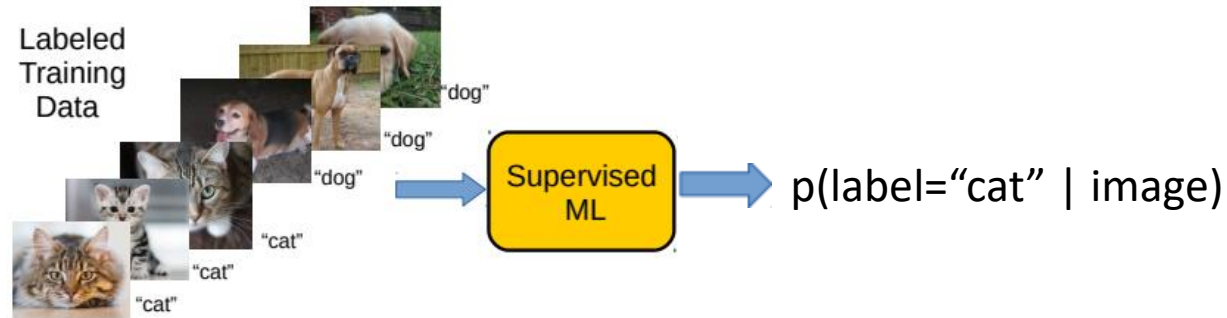


- Unsupervised Learning (“model/compress inputs”) can be thought of as estimating the **probability density** of the inputs



Perspective as probability estimation

- Supervised Learning (“predict output given input”) can be thought of as estimating the **conditional probability** of each possible output given an input



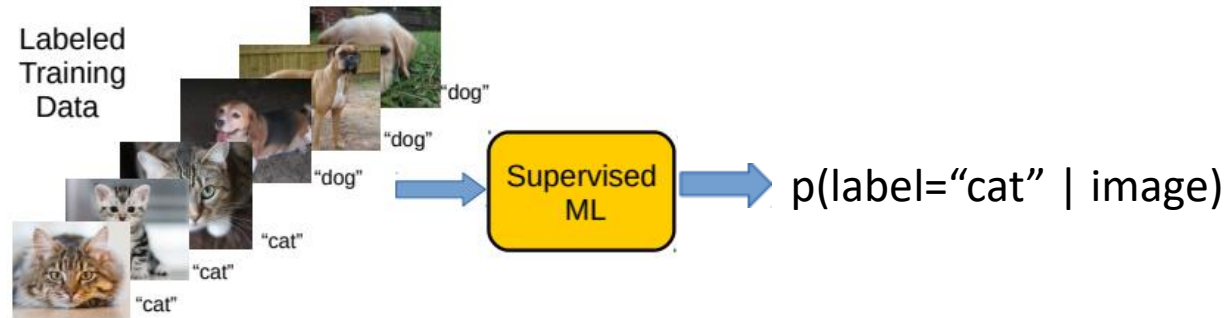
- Unsupervised Learning (“model/compress inputs”) can be thought of as estimating the **probability density** of the inputs

Harder since we don't know the labels in this case




Perspective as probability estimation

- Supervised Learning (“predict output given input”) can be thought of as estimating the **conditional probability** of each possible output given an input



- Unsupervised Learning (“model/compress inputs”) can be thought of as estimating the **probability density** of the inputs

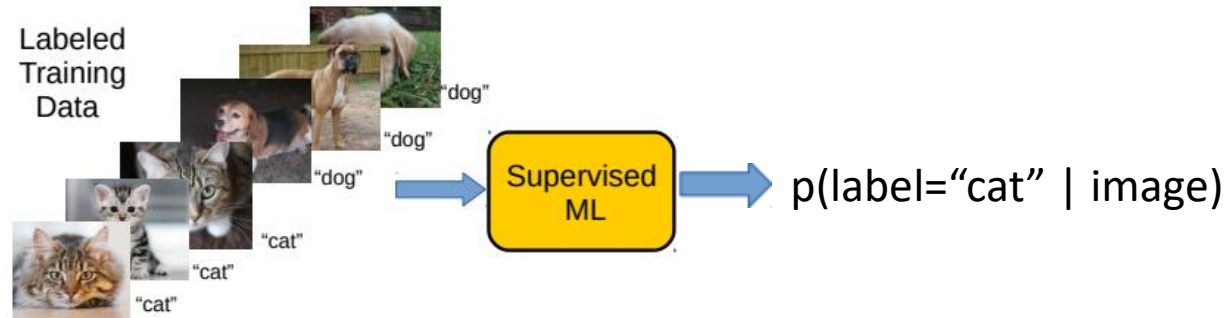


Harder since we don't know the labels in this case



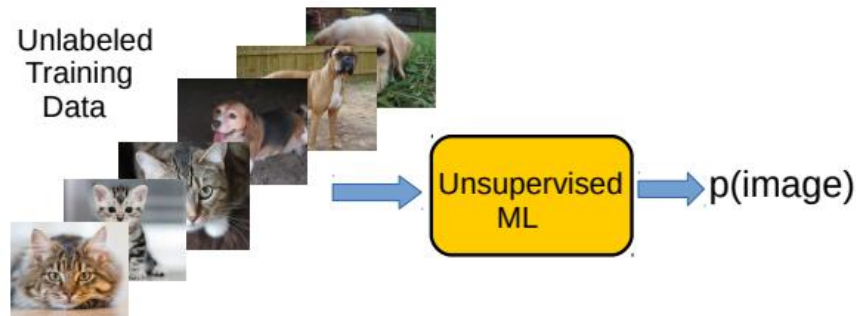
Perspective as probability estimation

- Supervised Learning (“predict output given input”) can be thought of as estimating the **conditional probability** of each possible output given an input



- Unsupervised Learning (“model/compress inputs”) can be thought of as estimating the **probability density** of the inputs

Harder since we don't know the labels in this case

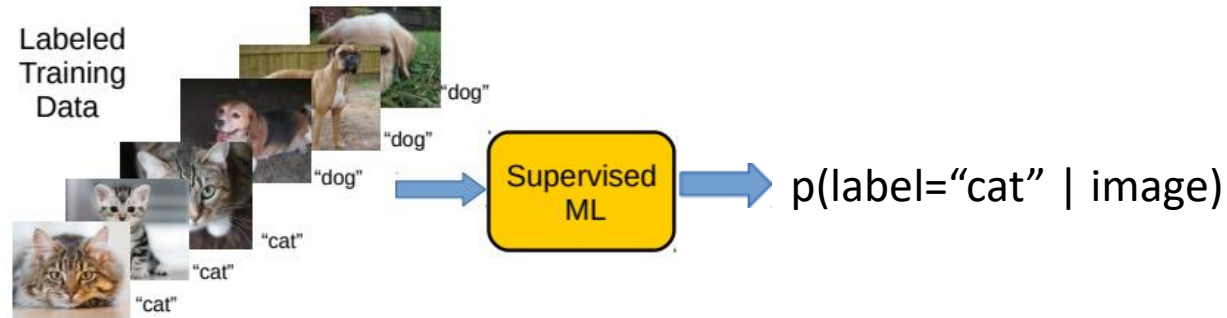


Don't worry if this doesn't make much sense as of now 😊 But the basic idea is to learn the underlying data distribution using the unlabeled inputs; many ways to do this as we will see later




Perspective as probability estimation

- Supervised Learning (“predict output given input”) can be thought of as estimating the **conditional probability** of each possible output given an input



- Unsupervised Learning (“model/compress inputs”) can be thought of as estimating the **probability density** of the inputs

Harder since we don't know the labels in this case



Don't worry if this doesn't make much sense as of now 😊 But the basic idea is to learn the underlying data distribution using the unlabeled inputs; many ways to do this as we will see later



- Reinforcement Learning can also be seen as estimating probability densities

Data and Features



Data and Features



Data and Features

- ML algos require a numeric **feature representation** of the inputs



Data and Features

- ML algos require a numeric **feature representation** of the inputs

Features represent semantics of the inputs. Being able to extract good features is key to the success of ML algos

14



Data and Features

Features represent semantics of the inputs. Being able to extract good features is key to the success of ML algos

14



- ML algos require a numeric **feature representation** of the inputs
- Features can be obtained using one of the two approaches



Data and Features

Features represent semantics of the inputs. Being able to extract good features is key to the success of ML algos

14



- ML algos require a numeric **feature representation** of the inputs
- Features can be obtained using one of the two approaches
 - Approach 1: Extracting/constructing features manually from raw inputs



Data and Features

Features represent semantics of the inputs. Being able to extract good features is key to the success of ML algos

14



- ML algos require a numeric **feature representation** of the inputs
- Features can be obtained using one of the two approaches
 - Approach 1: Extracting/constructing features manually from raw inputs
 - Approach 2: Learning the features from raw inputs



Data and Features

Features represent semantics of the inputs. Being able to extract good features is key to the success of ML algos

14



- ML algos require a numeric **feature representation** of the inputs
- Features can be obtained using one of the two approaches
 - Approach 1: Extracting/constructing features manually from raw inputs
 - Approach 2: Learning the features from raw inputs
- Approach 1 is what we will focus on primarily for now



Data and Features

Features represent semantics of the inputs. Being able to extract good features is key to the success of ML algos

14



- ML algos require a numeric **feature representation** of the inputs
- Features can be obtained using one of the two approaches
 - Approach 1: Extracting/constructing features manually from raw inputs
 - Approach 2: Learning the features from raw inputs
- Approach 1 is what we will focus on primarily for now
- Approach 2 is what is followed in **Deep Learning** algorithms (will see later)



Data and Features

Features represent semantics of the inputs. Being able to extract good features is key to the success of ML algos

14



- ML algos require a numeric **feature representation** of the inputs
- Features can be obtained using one of the two approaches
 - Approach 1: Extracting/constructing features manually from raw inputs
 - Approach 2: Learning the features from raw inputs
- Approach 1 is what we will focus on primarily for now
- Approach 2 is what is followed in **Deep Learning** algorithms (will see later)
- Approach 1 is not as powerful as Approach 2 but still used widely



Example: Feature Extraction for Text Data



Example: Feature Extraction for Text Data

- Consider some text data consisting of the following sentences:



Example: Feature Extraction for Text Data

- Consider some text data consisting of the following sentences:
 - John likes to watch movies



Example: Feature Extraction for Text Data

- Consider some text data consisting of the following sentences:
 - John likes to watch movies
 - Mary likes movies too



Example: Feature Extraction for Text Data

- Consider some text data consisting of the following sentences:
 - John likes to watch movies
 - Mary likes movies too
 - John also likes football



Example: Feature Extraction for Text Data

- Consider some text data consisting of the following sentences:
 - John likes to watch movies
 - Mary likes movies too
 - John also likes football
- Want to construct a **feature representation** for these sentences



Example: Feature Extraction for Text Data

- Consider some text data consisting of the following sentences:
 - John likes to watch movies
 - Mary likes movies too
 - John also likes football
- Want to construct a **feature representation** for these sentences
- Here is a “**bag-of-words**” (BoW) feature representation of these sentences



Example: Feature Extraction for Text Data

- Consider some text data consisting of the following sentences:
 - John likes to watch movies
 - Mary likes movies too
 - John also likes football
- Want to construct a **feature representation** for these sentences
- Here is a “**bag-of-words**” (BoW) feature representation of these sentences

	John	likes	to	watch	movies	Mary	too	also	football
Sentence 1	1	1	1	1	1	0	0	0	0
Sentence 2	0	1	0	0	1	1	1	0	0
Sentence 3	1	1	0	0	0	0	0	1	1



Example: Feature Extraction for Text Data

- Consider some text data consisting of the following sentences:

- John likes to watch movies
- Mary likes movies too
- John also likes football

BoW is just one of the many ways of doing feature extraction for text data. Not the most optimal one, and has various flaws (can you think of some?), but often works reasonably well



- Want to construct a **feature representation** for these sentences
- Here is a “**bag-of-words**” (BoW) feature representation of these sentences

	John	likes	to	watch	movies	Mary	too	also	football
Sentence 1	1	1	1	1	1	0	0	0	0
Sentence 2	0	1	0	0	1	1	1	0	0
Sentence 3	1	1	0	0	0	0	0	1	1



Example: Feature Extraction for Text Data

- Consider some text data consisting of the following sentences:

- John likes to watch movies
- Mary likes movies too
- John also likes football

BoW is just one of the many ways of doing feature extraction for text data. Not the most optimal one, and has various flaws (can you think of some?), but often works reasonably well



- Want to construct a **feature representation** for these sentences
- Here is a **“bag-of-words”** (BoW) feature representation of these sentences

	John	likes	to	watch	movies	Mary	too	also	football
Sentence 1	1	1	1	1	1	0	0	0	0
Sentence 2	0	1	0	0	1	1	1	0	0
Sentence 3	1	1	0	0	0	0	0	1	1

- Each sentence is now represented as a **binary vector** (each feature is a binary value, denoting presence or absence of a word). BoW is also called **“unigram”** rep.



Example: Feature Extraction for Image Data




Example: Feature Extraction for Image Data

- A very simple feature extraction approach for image data is **flattening**




Example: Feature Extraction for Image Data

- A very simple feature extraction approach for image data is **flattening**



Example: Feature Extraction for Image Data

- A very simple feature extraction approach for image data is **flattening**




- **Histogram** of visual patterns is another popular feature extr. method for images




Example: Feature Extraction for Image Data

- A very simple feature extraction approach for image data is **flattening**




- **Histogram** of visual patterns is another popular feature extr. method for images



Example: Feature Extraction for Image Data


- A very simple feature extraction approach for image data is **flattening**



Flattening and histogram based methods destroy the spatial information in the image but often still work reasonably well




- **Histogram** of visual patterns is another popular feature extr. method for images



Example: Feature Extraction for Image Data


- A very simple feature extraction approach for image data is **flattening**



Flattening and histogram based methods destroy the spatial information in the image but often still work reasonably well



- **Histogram** of visual patterns is another popular feature extr. method for images



- Many other manual feature extraction techniques developed in computer vision and image processing communities (SIFT, HoG, and others)



Feature Selection

17



Feature Selection

- Not all the extracted features may be relevant for learning the model (some may even confuse the learner)




Feature Selection

- Not all the extracted features may be relevant for learning the model (some may even confuse the learner)
- **Feature selection** (a step after feature extraction) can be used to identify the features that matter, and discard the others, for more effective learning




Feature Selection

- Not all the extracted features may be relevant for learning the model (some may even confuse the learner)
- **Feature selection** (a step after feature extraction) can be used to identify the features that matter, and discard the others, for more effective learning




Feature Selection

- Not all the extracted features may be relevant for learning the model (some may even confuse the learner)
- **Feature selection** (a step after feature extraction) can be used to identify the features that matter, and discard the others, for more effective learning



Feature Selection

- Not all the extracted features may be relevant for learning the model (some may even confuse the learner)
- **Feature selection** (a step after feature extraction) can be used to identify the features that matter, and discard the others, for more effective learning



Calculating BMI from this data doesn't require ML but this simple example is just to illustrate the idea of feature selection 😊



Feature Selection

- Not all the extracted features may be relevant for learning the model (some may even confuse the learner)
- **Feature selection** (a step after feature extraction) can be used to identify the features that matter, and discard the others, for more effective learning

Age
Gender
Height
Weight
Eye color



Body-mass index (BMI)

Calculating BMI from this data doesn't require ML but this simple example is just to illustrate the idea of feature selection 😊




- Many techniques exist – some based on intuition, some based on algorithmic principles (will visit feature selection later)



Feature Selection

- Not all the extracted features may be relevant for learning the model (some may even confuse the learner)
- **Feature selection** (a step after feature extraction) can be used to identify the features that matter, and discard the others, for more effective learning



Calculating BMI from this data doesn't require ML but this simple example is just to illustrate the idea of feature selection 😊



- Many techniques exist – some based on intuition, some based on algorithmic principles (will visit feature selection later)
- More common in supervised learning but can also be done for unsup. learning



Some More Postprocessing: Feature Scaling



Some More Postprocessing: Feature Scaling

- Even after feature selection, the features may not be on the same scale



Some More Postprocessing: Feature Scaling

- Even after feature selection, the features may not be on the same scale
- This can be problematic when comparing two inputs – features that have larger scales may dominate the result of such comparisons



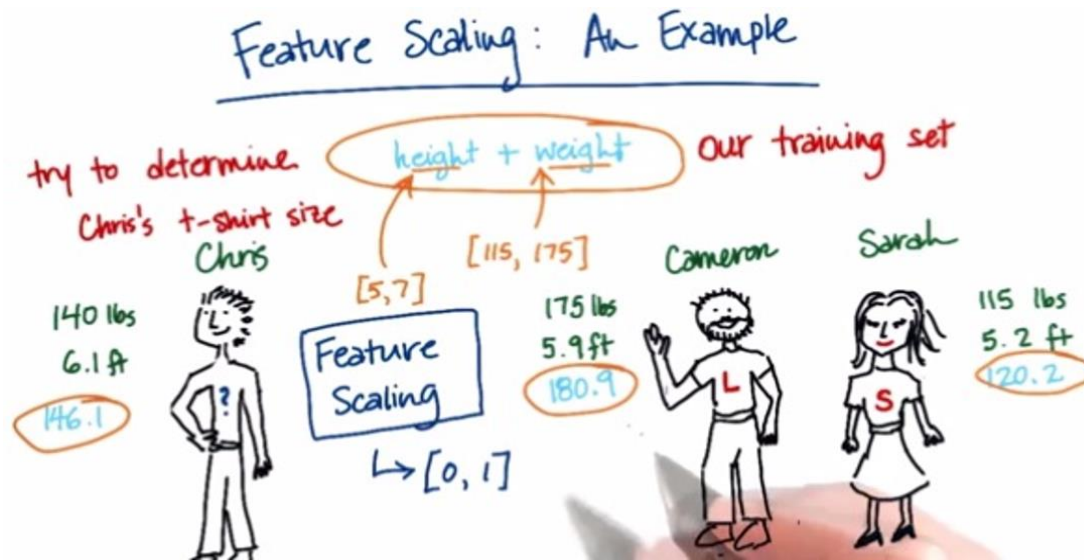
Some More Postprocessing: Feature Scaling

- Even after feature selection, the features may not be on the same scale
- This can be problematic when comparing two inputs – features that have larger scales may dominate the result of such comparisons
- Therefore helpful to standardize the features (e.g., by bringing all of them on the same scale such as between 0 to 1)



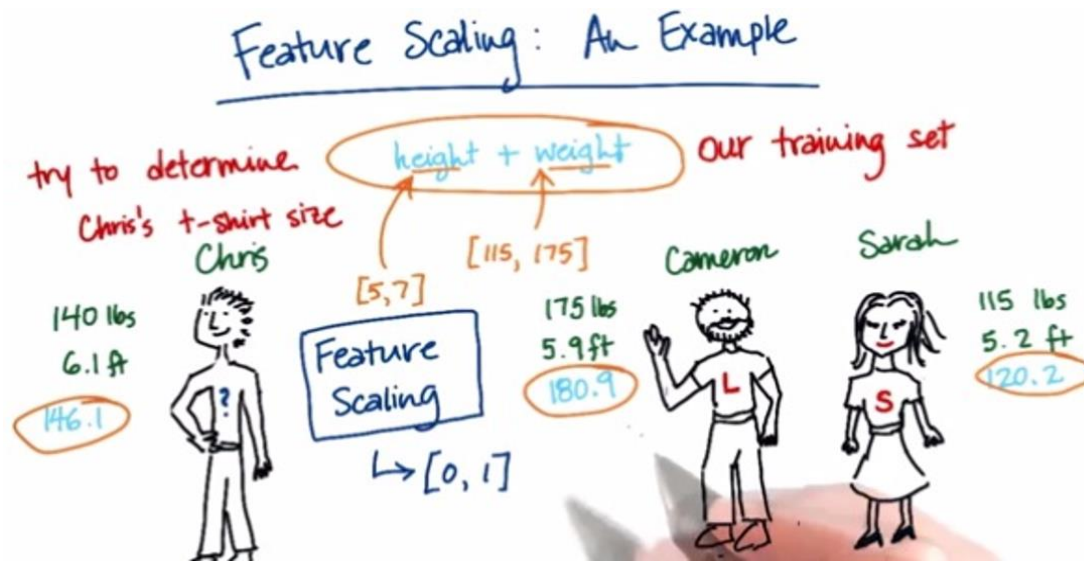
Some More Postprocessing: Feature Scaling

- Even after feature selection, the features may not be on the same scale
- This can be problematic when comparing two inputs – features that have larger scales may dominate the result of such comparisons
- Therefore helpful to standardize the features (e.g., by bringing all of them on the same scale such as between 0 to 1)



Some More Postprocessing: Feature Scaling

- Even after feature selection, the features may not be on the same scale
- This can be problematic when comparing two inputs – features that have larger scales may dominate the result of such comparisons
- Therefore helpful to standardize the features (e.g., by bringing all of them on the same scale such as between 0 to 1)






- Also helpful for stabilizing the optimization techniques used in ML algos



Some More Postprocessing: Feature Scaling

- Even after feature selection, the features may not be on the same scale
- This can be problematic when comparing two inputs – features that have larger scales may dominate the result of such comparisons
- Therefore helpful to standardize the features (e.g., by bringing all of them on the same scale such as between 0 to 1)



- Also helpful for stabilizing the optimization techniques used in ML algos



Deep Learning: An End-to-End Approach to ML

19



Deep Learning: An End-to-End Approach to ML

Deep Learning = ML with **automated feature learning** from the raw inputs



Deep Learning: An End-to-End Approach to ML

Deep Learning = ML with **automated feature learning** from the raw inputs


Feature extraction part is automated via the feature learning module



Deep Learning: An End-to-End Approach to ML

Deep Learning = ML with **automated feature learning** from the raw inputs


Feature extraction part is automated via the feature learning module



Deep Learning: An End-to-End Approach to ML

Deep Learning = ML with **automated feature learning** from the raw inputs

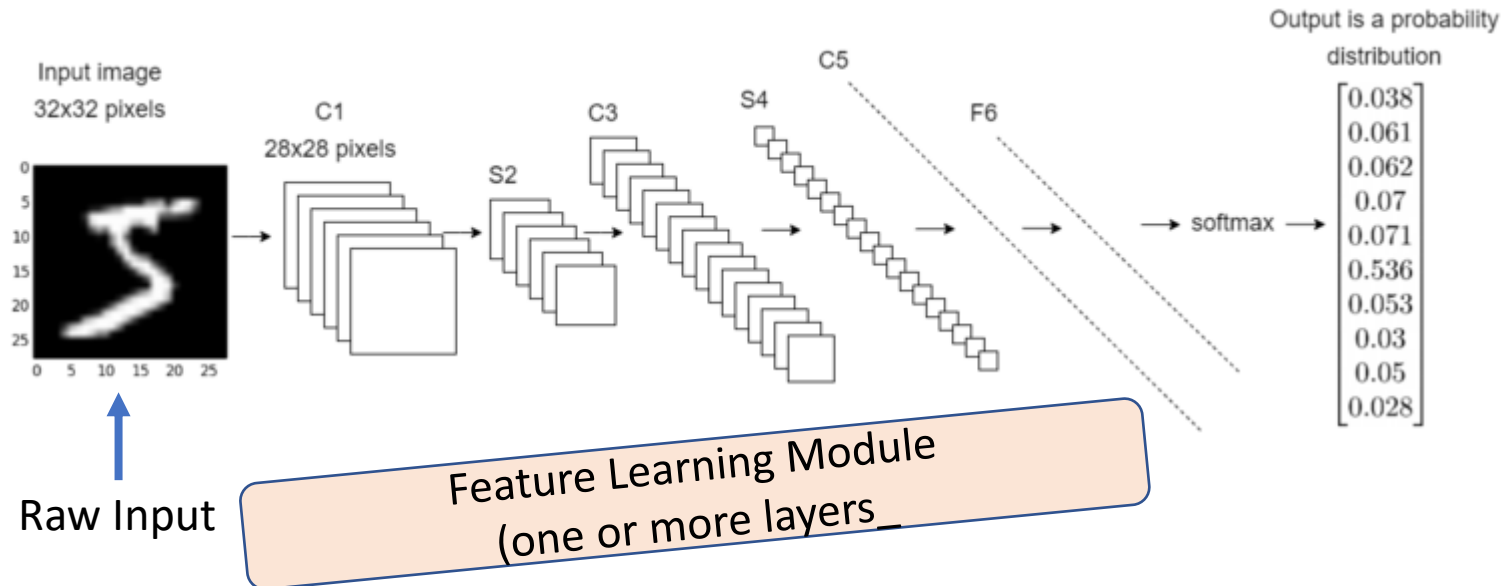
Feature extraction part is automated via the feature learning module



Deep Learning: An End-to-End Approach to ML

Deep Learning = ML with **automated feature learning** from the raw inputs

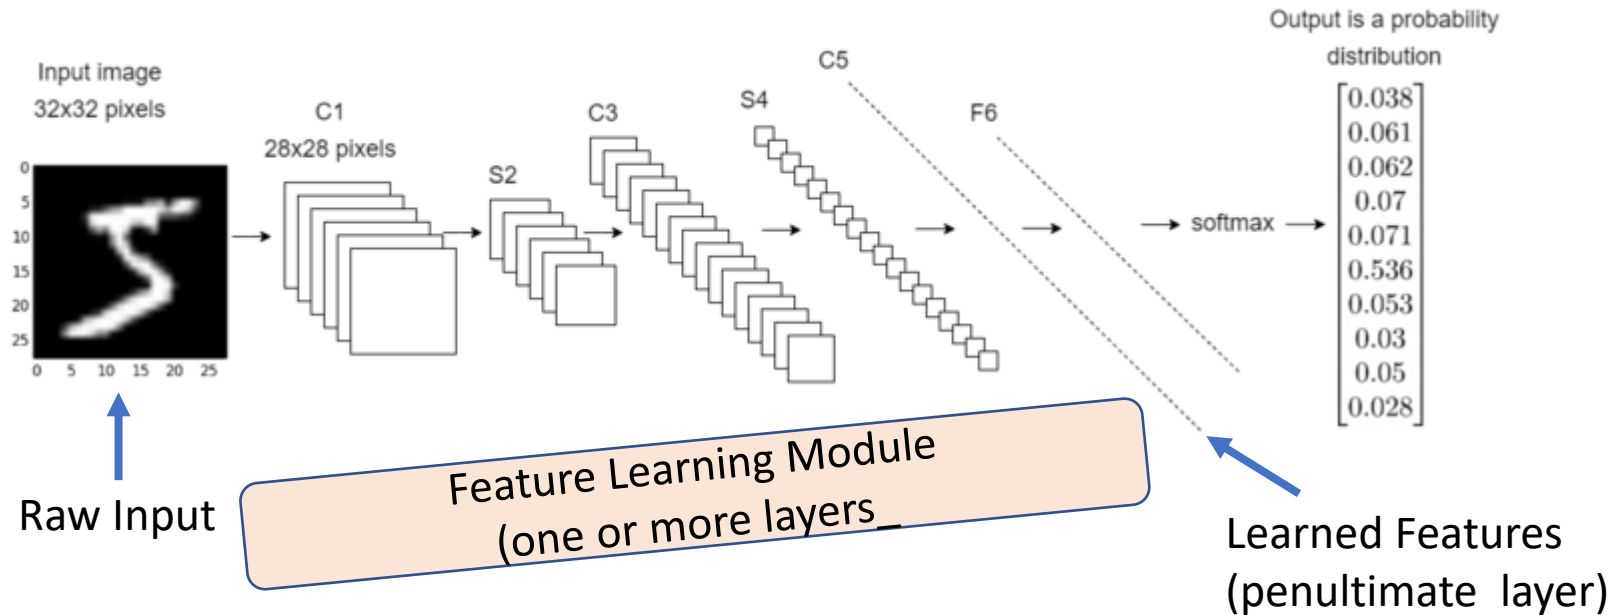
Feature extraction part is automated via the feature learning module



Deep Learning: An End-to-End Approach to ML

Deep Learning = ML with **automated feature learning** from the raw inputs

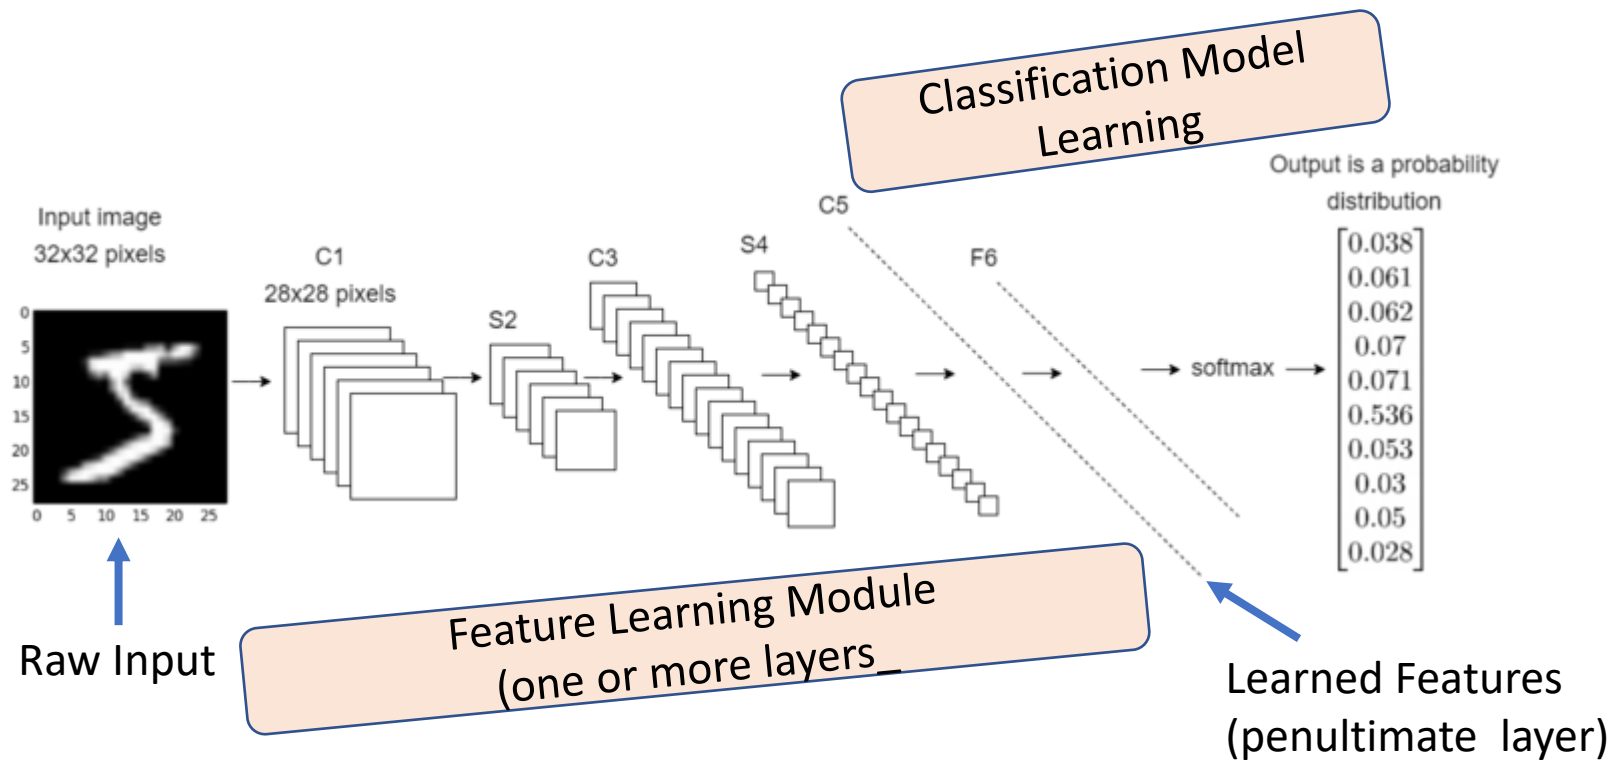
Feature extraction part is automated via the feature learning module



Deep Learning: An End-to-End Approach to ML

Deep Learning = ML with **automated feature learning** from the raw inputs


Feature extraction part is automated via the feature learning module



Deep Learning: An End-to-End Approach to ML

Deep Learning = ML with **automated feature learning** from the raw inputs

Feature extraction part is automated via the feature learning module



Some Notation/Nomenclature/Convention

input-output pairs $\{(\mathbf{x}_n, y_n)\}_{n=1}^N$

- inputs $\{\mathbf{x}_n\}_{n=1}^N$
- is (usually) a vector containing the values of the **features** or **attributes** or **covariates** that encode properties of the it represents, e.g.,
 - can be a 49×1 vector of pixel intensities
- is the **output** or **response** or **label** associated with input \mathbf{x}_n (and its value is known for the training inputs)



Some Notation/Nomenclature/Convention

- \mathbf{x}_n, y_n \mathbf{x}_n \mathbf{x} \mathbf{x}_n \mathbf{n} \mathbf{x}_n, y_n y y_n n y_n \mathbf{x}_n, y_n $\{ \mathbf{x}_n, y_n \}_{n=1}^N$
 $n=1 \{ \mathbf{x}_n, y_n \}_{n=1}^N$ $NN \{ \mathbf{x}_n, y_n \}_{n=1}^N$
- *Sup.* learning requires training data as N input-output pairs $\{ \mathbf{x}_n, y_n \}_{n=1}^N$

- inputs $\{\mathbf{x}_n\}_{n=1}^N$

- is (usually) a vector containing the values of the **features** or **attributes** or **covariates** that encode properties of the it represents, e.g.,

- can be a 49×1 vector of pixel intensities

- is the **output** or **response** or **label** associated with input \mathbf{x}_n (and its value is known for the training inputs)



Some Notation/Nomenclature/Convention

- \mathbf{x}_n $\mathbf{x}_1 \mathbf{x}_2 \dots \mathbf{x}_N$ $\{\mathbf{x}_n\}_{n=1}^N$ \mathbf{X} \mathbf{x}_n $\{ \mathbf{x}_n \}_{n=1}^N$ N N $\{ \mathbf{x}_n \}_{n=1}^N$ N
- y_n $y_1 y_2 \dots y_N$ $\{y_n\}_{n=1}^N$ \mathbf{y} $\{ \mathbf{y}_n \}_{n=1}^N$ N N $\{ \mathbf{y}_n \}_{n=1}^N$ N
- *Sup.* learning requires training data as N input-output pairs $\{ \mathbf{x}_n, y_n \}_{n=1}^N$
- Unsupervised learning requires training data as N inputs $\{ \mathbf{x}_n \}_{n=1}^N$
- inputs $\{\mathbf{x}_n\}_{n=1}^N$
- \mathbf{x}_n is (usually) a vector containing the values of the **features** or **attributes** or **covariates** that encode properties of the it represents, e.g.,
 - \mathbf{x}_n can be a 49×1 vector of pixel intensities
- y_n is the **output** or **response** or **label** associated with input \mathbf{x}_n (and its value is known for the training inputs)



Some Notation/Nomenclature/Convention

- $\{\mathbf{x}_n\}_{n=1}^N$
- $\{\mathbf{x}_n, y_n\}_{n=1}^N$
- *Sup.* learning requires training data as N input-output pairs $\{\mathbf{x}_n, y_n\}_{n=1}^N$
- Unsupervised learning requires training data as N inputs $\{\mathbf{x}_n\}_{n=1}^N$
- inputs $\{\mathbf{x}_n\}_{n=1}^N$

RL and other flavors of ML problems also use similar notation



- \mathbf{x}_n is (usually) a vector containing the values of the **features** or **attributes** or **covariates** that encode properties of the it represents, e.g.,

- can be a 49×1 vector of pixel intensities

- y_n is the **output** or **response** or **label** associated with input \mathbf{x}_n (and its value is known for the training inputs)



Some Notation/Nomenclature/Convention

- is (usually) a vector containing the values of the features or attributes or covariates that encode properties of the it represents, e.g.,

- $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N$ or $\{\mathbf{x}_n\}_{n=1}^N$

RL and other flavors of ML problems also use similar notation



- y_1, y_2, \dots, y_N or $\{y_n\}_{n=1}^N$

- *Sup.* learning requires training data as N input-output pairs $\{\mathbf{x}_n, y_n\}_{n=1}^N$

- Unsupervised learning requires training data as N inputs $\{\mathbf{x}_n\}_{n=1}^N$

- Each input \mathbf{x}_n is (usually) a vector containing the values of the **features** or **attributes** or **covariates** that encode properties of the it represents, e.g.,

- is (usually) a vector containing the values of the **features** or **attributes** or **covariates** that encode properties of the it represents, e.g.,

- can be a 49×1 vector of pixel intensities



Some Notation/Nomenclature/Convention

- can be a 49×1 vector of pixel intensities
- is (usually) a vector containing the values of the features or attributes or covariates that encode properties of the it represents, e.g.,
- $\{ \mathbf{x}_n \}_{n=1}^N$
- $\{ \mathbf{x}_n, y_n \}_{n=1}^N$
- *Sup.* learning requires training data as N input-output pairs $\{ \mathbf{x}_n, y_n \}_{n=1}^N$
- Unsupervised learning requires training data as N input
- Each input \mathbf{x}_n is (usually) a vector containing the values of the **features** or **attributes** or **covariates** that encode properties of the it represents, e.g.,
 - For a 7×7 image: \mathbf{x}_n can be a 49×1 vector of pixel intensities
- is (usually) a vector containing the values of the **features** or **attributes** or **covariates** that encode properties of the it represents, e.g.,

RL and other flavors of ML problems also use similar notation



Size or length of the input \mathbf{x}_n is commonly known as **data/input dimensionality** or **feature dimensionality**



Some Notation/Nomenclature/Convention

- is the **output** or **response** or **label** associated with input \mathbf{x}_n (and its value is known for the training inputs)
- can be a 49×1 vector of pixel intensities
- is (usually) a vector containing the values of the features or attributes or covariates that encode properties of the it represents, e.g.,
 - $\{\mathbf{x}_n\}_{n=1}^N$
 - $\{\mathbf{x}_n, y_n\}_{n=1}^N$
- *Sup.* learning requires training data as N input-output pairs $\{\mathbf{x}_n, y_n\}_{n=1}^N$
- Unsupervised learning requires training data as N inputs $\{\mathbf{x}_n\}_{n=1}^N$
- Each input \mathbf{x}_n is (usually) a vector containing the values of the **features** or **attributes** or **covariates** that encode properties of the it represents, e.g.,
 - For a 7×7 image: \mathbf{x}_n can be a 49×1 vector of pixel intensities

RL and other flavors of ML problems also use similar notation



Size or length of the input \mathbf{x}_n is commonly known as **data/input dimensionality** or **feature dimensionality**



Some Notation/Nomenclature/Convention

- is the **output** or **response** or **label** associated with input \mathbf{x}_n (and its value is known for the training inputs)
- can be a 49×1 vector of pixel intensities
- is (usually) a vector containing the values of the features or attributes or covariates that encode properties of the it represents, e.g.,
 - $\{\mathbf{x}_n\}_{n=1}^N$
 - $\{\mathbf{x}_n, y_n\}_{n=1}^N$
- *Sup.* learning requires training data as N input-output pairs $\{\mathbf{x}_n, y_n\}_{n=1}^N$
- Unsupervised learning requires training data as N inputs $\{\mathbf{x}_n\}_{n=1}^N$
- Each input \mathbf{x}_n is (usually) a vector containing the values of the **features** or **attributes** or **covariates** that encode properties of the it represents, e.g.,
 - For a 7×7 image: \mathbf{x}_n can be a 49×1 vector of pixel intensities

RL and other flavors of ML problems also use similar notation



Size or length of the input \mathbf{x}_n is commonly known as **data/input dimensionality** or **feature dimensionality**



Types of Features and Types of Outputs



Types of Features and Types of Outputs

- Features as well as outputs can be real-valued, binary, categorical, ordinal, etc.



Types of Features and Types of Outputs

- Features as well as outputs can be real-valued, binary, categorical, ordinal, etc.
- **Real-valued:** Pixel intensity, house area, house price, rainfall amount, temperature, etc



Types of Features and Types of Outputs

- Features as well as outputs can be real-valued, binary, categorical, ordinal, etc.
- **Real-valued:** Pixel intensity, house area, house price, rainfall amount, temperature, etc
- **Binary:** Male/female, adult/non-adult, or any yes/no or present/absent type value



Types of Features and Types of Outputs

- Features as well as outputs can be real-valued, binary, categorical, ordinal, etc.
- **Real-valued:** Pixel intensity, house area, house price, rainfall amount, temperature, etc
- **Binary:** Male/female, adult/non-adult, or any yes/no or present/absent type value
- **Categorical/Discrete:** Zipcode, blood-group, or any “one from a finite many choices” value




Types of Features and Types of Outputs

- Features as well as outputs can be real-valued, binary, categorical, ordinal, etc.
- **Real-valued:** Pixel intensity, house area, house price, rainfall amount, temperature, etc
- **Binary:** Male/female, adult/non-adult, or any yes/no or present/absent type value
- **Categorical/Discrete:** Zipcode, blood-group, or any “one from a finite many choices” value
- **Ordinal:** Grade (A/B/C etc.) in a course, or any other type where relative values matter



Types of Features and Types of Outputs

- Features as well as outputs can be real-valued, binary, categorical, ordinal, etc.
- **Real-valued:** Pixel intensity, house area, house price, rainfall amount, temperature, etc
- **Binary:** Male/female, adult/non-adult, or any yes/no or present/absent type value
- **Categorical/Discrete:** Zipcode, blood-group, or any “one from a finite many choices” value
- **Ordinal:** Grade (A/B/C etc.) in a course, or any other type where relative values matter
- Often, the features can be of mixed types (some real, some categorical, some ordinal, etc.)



Some Basic Operations of Inputs

$\in R^D$ to of size D

- inputs $\{\mathbf{x}_n\}_{n=1}^N$, their average or mean can be computed as

$$= \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$$

- and \mathbf{x}_m

- and the mean $\boldsymbol{\mu}$ of all inputs



Some Basic Operations of Inputs

- $R D R R R D D D R D$ to of size D
- Assume each input feature vector $x_n \in \mathbb{R}^D$ to of size D
- inputs $\{\mathbf{x}_n\}_{n=1}^N$, their average or mean can be computed as

$$= \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$$
- and \mathbf{x}_m
- and the mean $\boldsymbol{\mu}$ of all inputs



Some Basic Operations of Inputs

- $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N$ their average or mean can be computed as
- $\mathbf{x}_n \in \mathbb{R}^D$ to of size D
- Assume each input feature vector $\mathbf{x}_n \in \mathbb{R}^D$ to of size D
- Given N inputs $\{\mathbf{x}_n\}_{n=1}^N$ their average or mean can be computed as
- inputs $\{\mathbf{x}_n\}_{n=1}^N$, their average or mean can be computed as

$$= \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$$

- and \mathbf{x}_m

- and the mean μ of all inputs



Some Basic Operations of Inputs

- $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N$ are N input feature vectors of size D
- $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N\}$ their average or mean can be computed as
- $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N$ to of size D
- Assume each input feature vector $\mathbf{x}_n \in \mathbb{R}^D$ to of size D

$$\mu = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$$

- inputs $\{\mathbf{x}_n\}_{n=1}^N$, their average or mean can be computed as
- $$= \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$$

- and \mathbf{x}_m



Some Basic Operations of Inputs

- $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N$ are input feature vectors of size D
- $\{\mathbf{x}_n\}_{n=1}^N$ is a set of N input feature vectors of size D
- $\bar{\mathbf{x}} = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$ is the mean of the input feature vectors
- $\bar{\mathbf{x}}$ is the average or mean of the input feature vectors
- Assume each input feature vector $\mathbf{x}_n \in \mathbb{R}^D$ to of size D

What does such a
"mean" represent?



- inputs $\{\mathbf{x}_n\}_{n=1}^N$, their average or mean can be computed as

$$\bar{\mathbf{x}} = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$$

- and \mathbf{x}_m



Some Basic Operations of Inputs

- $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N$ are input feature vectors of size D

- The mean of the input feature vectors $\{\mathbf{x}_n\}_{n=1}^N$ can be computed as

- The mean vector μ is of size D

- Assume each input feature vector $\mathbf{x}_n \in \mathbb{R}^D$ to of size D

$$\mu = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$$


- For a set of inputs $\{\mathbf{x}_n\}_{n=1}^N$, their average or mean can be computed as

$$= \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$$

- and \mathbf{x}_m

What does such a "mean" represent?

If inputs are all cat images, mean vector would represent what an "average" cat looks like



Some Basic Operations of Inputs

- and \mathbf{x}_m
- $\frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$ What does such a "mean" represent?
- $\{\mathbf{x}_n\}_{n=1}^N$ their average mean can be computed as
- $\mathbf{x}_n \in \mathbb{R}^D$ to of size D
- Assume each input feature vector $\mathbf{x}_n \in \mathbb{R}^D$ to of size D

$$\mu = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$$

- Can compute the Euclidean distance between any pair of inputs \mathbf{x}_n and \mathbf{x}_m

$$= \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$$

- and \mathbf{x}_m

If inputs are all cat images, mean vector would represent what an "average" cat looks like



Some Basic Operations of Inputs

- and $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N$
- $\frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$ What does such a "mean" represent?
- $\{\mathbf{x}_n\}_{n=1}^N$ their average mean can be computed as
- \mathbf{R} to of size D
- Assume each input feature vector $\mathbf{x}_n \in \mathbb{R}^D$ to of size D

If inputs are all cat images, mean vector would represent what an "average" cat looks like



- Can compute distance between \mathbf{x}_n and \mathbf{x}_m as

$$d(\mathbf{x}_n, \mathbf{x}_m) = \|\mathbf{x}_n - \mathbf{x}_m\| = \sqrt{(\mathbf{x}_n - \mathbf{x}_m)^\top (\mathbf{x}_n - \mathbf{x}_m)} = \sqrt{\sum_{d=1}^D (x_{nd} - x_{md})^2}$$

$$= \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$$

and \mathbf{x}_m



Some Basic Operations of Inputs

- and the mean μ of all inputs

- and \mathbf{x}_m

- $\frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$

- $\{\mathbf{x}_n\}_{n=1}^N$ mean can be computed as

- \mathbf{R}^D to of size D

- Assum

$$d(\mathbf{x}_n, \mathbf{x}_m) = \|\mathbf{x}_n - \mathbf{x}_m\| = \sqrt{(\mathbf{x}_n - \mathbf{x}_m)^\top (\mathbf{x}_n - \mathbf{x}_m)} = \sqrt{\sum_{d=1}^D (x_{nd} - x_{md})^2}$$

- Can compute the Euclidean distance between any pair of inputs \mathbf{x}_n and \mathbf{x}_m

$$= \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$$

What does such a "mean" represent?

If inputs are all cat images, mean vector would represent what an "average" cat looks like



- or Euclidean distance between an input \mathbf{x}_n and the mean μ of all inputs

Some Basic Operations of Inputs

- and the mean μ of all inputs

- and \mathbf{x}_m

- $\frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$

- $\{\mathbf{x}_n\}_{n=1}^N$ mean can be computed as

- \mathbf{R}^D to of size D

- Assum

$$d(\mathbf{x}_n, \mathbf{x}_m) = \|\mathbf{x}_n - \mathbf{x}_m\| = \sqrt{(\mathbf{x}_n - \mathbf{x}_m)^\top (\mathbf{x}_n - \mathbf{x}_m)} = \sqrt{\sum_{d=1}^D (x_{nd} - x_{md})^2}$$

- Can compute the Euclidean distance between any pair of inputs \mathbf{x}_n and \mathbf{x}_m

$$= \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n$$

What does such a "mean" represent?

If inputs are all cat images, mean vector would represent what an "average" cat looks like



- or Euclidean distance between an input \mathbf{x}_n and the mean μ of all inputs

Next Class

- Introduction to Supervised Learning
- A simple Supervised Learning algorithm based on computing distances

