

# Turning Linear Models into Nonlinear Models using Kernel Methods (Contd)

CS771: Introduction to Machine Learning

Piyush Rai

# Using Kernels

- Kernels can turn many linear models into nonlinear models
- Recall that  $k(\mathbf{x}, \mathbf{z})$  represents a dot product in some high-dim feature space  $\mathcal{F}$
- Important: Any ML model/algo in which, during training and test, inputs only appear as dot product can be “kernelized”
- Just replace each term of the form  $\mathbf{x}_i^\top \mathbf{x}_j$  by  $\phi(\mathbf{x}_i)^\top \phi(\mathbf{x}_j) = k(\mathbf{x}_i, \mathbf{x}_j) = K_{ij}$
- Most ML models/algos can be easily kernelized, e.g.,
  - Distance based methods, Perceptron, SVM, linear regression, etc.
  - Many of the unsupervised learning algorithms too can be kernelized (e.g., K-means clustering, Principal Component Analysis, etc. - will see later)
  - Let’s look at two examples: Kernelized SVM and Kernelized Ridge Regression



# An Aside: Kernelizing a Euclidean Distance

- Many algorithms, e.g., LwP, KNN, etc. use Euclidean distances, e.g.,

$$d(a, b) = \|a - b\|^2 = \|a\|^2 + \|b\|^2 - 2a^\top b = a^\top a + b^\top b - 2a^\top b$$

- This can be kernelized as well by replacing the above norms and inner products by their kernelized versions, assuming a kernel  $k$  with feature map  $\phi$

$$\begin{aligned} d(\phi(a), \phi(b)) &= \|\phi(a) - \phi(b)\|^2 \\ &= \phi(a)^\top \phi(a) + \phi(b)^\top \phi(b) - 2\phi(a)^\top \phi(b) \\ &= k(a, a) + k(b, b) - 2k(a, b) \end{aligned}$$



# Nonlinear SVM using Kernels



# Kernelized SVM Training


- Recall the soft-margin linear SVM objective (with no bias term)

$$\operatorname{argmax}_{\mathbf{0} \leq \boldsymbol{\alpha} \leq \mathbf{C}} \quad \boldsymbol{\alpha}^T \mathbf{1} - \frac{1}{2} \boldsymbol{\alpha}^T \mathbf{G} \boldsymbol{\alpha}$$

Inputs only appear  
as dot products ☺

$$G_{ij} = y_i y_j \mathbf{x}_i^T \mathbf{x}_j$$

- To kernelize, we can simply replace  $G_{ij} = y_i y_j \mathbf{x}_i^T \mathbf{x}_j$  by  $y_i y_j K_{ij}$ 
  - .. where  $K_{ij} = k(\mathbf{x}_i, \mathbf{x}_j) = \phi(\mathbf{x}_i)^T \phi(\mathbf{x}_j)$  for a suitable kernel function  $k$
- The problem can now be solved just like the linear SVM case
- The new SVM learns a linear separator in kernel-induced feature space  $\mathcal{F}$ 
  - This corresponds to a **non-linear separator** in the original feature space  $\mathcal{X}$



# Kernelized SVM Prediction

- SVM weight vector for the kernelized case will be  $\mathbf{w} = \sum_{n=1}^N \alpha_n y_n \phi(\mathbf{x}_n)$
- Note: We can't store  $\mathbf{w}$  unless the feature mapping  $\phi(\mathbf{x}_n)$  is finite dimensional
  - In practice, we store the  $\alpha_n$ 's and the training data for test time (just like KNN)
  - In fact, need to store only training examples for which  $\alpha_n$  is nonzero (i.e., the support vectors)
- Prediction for a new test input  $\mathbf{x}_*$  (assuming hyperplane's bias  $b = 0$ ) will be
 
$$y_* = \text{sign}(\mathbf{w}^\top \phi(\mathbf{x}_*)) = \text{sign}\left(\sum_{n=1}^N \alpha_n y_n \phi(\mathbf{x}_n)^\top \phi(\mathbf{x}_*)\right) = \text{sign}\left(\sum_{n=1}^N \alpha_n y_n k(\mathbf{x}_n, \mathbf{x}_*)\right)$$
- Note that the prediction cost also scales linearly with  $N$  (unlike a linear model where we only need to compute  $\mathbf{w}^\top \mathbf{x}_*$ , whose cost only depends on  $D$ , not  $N$ )
- Also note that, for unkernelized (i.e., linear) SVM,  $\mathbf{w} = \sum_{n=1}^N \alpha_n y_n \mathbf{x}_n$  can be computed and stored as a  $D \times 1$  vector and we can compute  $\mathbf{w}^\top \mathbf{x}_*$  in  $O(D)$  time



# Nonlinear Ridge Regression using Kernels



# Kernelized Ridge Regression

- Recall the ridge regression problem:  $\mathbf{w} = \arg \min_{\mathbf{w}} \sum_{n=1}^N (y_n - \mathbf{w}^\top \mathbf{x}_n)^2 + \lambda \mathbf{w}^\top \mathbf{w}$

- The solution to this problem was

Inputs don't appear to be as inner product. No hope of kernelization?



They do; with a bit of algebra ☺

$$\mathbf{w} = \left( \sum_{n=1}^N \mathbf{x}_n \mathbf{x}_n^\top + \lambda \mathbf{I}_D \right) \left( \sum_{n=1}^N y_n \mathbf{x}_n \right) = (\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I}_D)^{-1} \mathbf{X}^\top \mathbf{y}$$

- Can use matrix inversion lemma  $(\mathbf{F}\mathbf{H}^{-1}\mathbf{G} - \mathbf{E})^{-1}\mathbf{F}\mathbf{H}^{-1} = \mathbf{E}^{-1}\mathbf{F}(\mathbf{G}\mathbf{E}^{-1}\mathbf{F} - \mathbf{H})^{-1}$

- Using the lemma, can rewrite  $\mathbf{w}$  as

$$\mathbf{w} = \mathbf{X}^\top (\mathbf{X}\mathbf{X}^\top + \lambda \mathbf{I}_N)^{-1} \mathbf{y} = \mathbf{X}^\top \boldsymbol{\alpha} = \sum_{n=1}^N \alpha_n \mathbf{x}_n \quad \text{where} \quad \boldsymbol{\alpha} = (\mathbf{X}\mathbf{X}^\top + \lambda \mathbf{I}_N)^{-1} \mathbf{y} = (\mathbf{K} + \lambda \mathbf{I}_N)^{-1} \mathbf{y}$$

$N \times 1$  vector of dual variables

Note: Not sparse unlike SVM

- Kernelized weight vector will be  $\mathbf{w} = \sum_{n=1}^N \alpha_n \phi(\mathbf{x}_n)$

Prediction cost is also linear in  $N$  (like KNN)

- Prediction for a test input  $\mathbf{x}_*$  will be  $\mathbf{w}^\top \phi(\mathbf{x}_*) = \sum_{n=1}^N \alpha_n \phi(\mathbf{x}_n)^\top \phi(\mathbf{x}_*) = \sum_{n=1}^N \alpha_n k(\mathbf{x}_n, \mathbf{x}_*)$





# Speeding-up Kernel Methods



# Speeding-up Kernel Methods

- Kernel methods, unlike linear models are slow at training and test time
- Would be nice if we could easily compute mapping  $\phi(\mathbf{x})$  associated with kernel  $k$
- Then we could apply linear models directly on  $\phi(\mathbf{x})$  without having to kernelize
- But this is in general not possible since  $\phi(\mathbf{x})$  is very high/infinite dimensional
- An alternative: Get a good set of low-dim features  $\psi(\mathbf{x}) \in \mathbb{R}^L$  using the kernel  $k$
- If  $\psi(\mathbf{x})$  is a good approximation to  $\phi(\mathbf{x})$  then we can use  $\psi(\mathbf{x})$  in a linear model

Goodness Criterion: 
$$\psi(\mathbf{x}_i)^\top \psi(\mathbf{x}_j) \approx \phi(\mathbf{x}_i)^\top \phi(\mathbf{x}_j)$$


... which also means  $\psi(\mathbf{x}_i)^\top \psi(\mathbf{x}_j) \approx k(\mathbf{x}_i, \mathbf{x}_j)$

- Will look at two popular approaches: [Landmarks](#) and [Random Features](#)



# Extracting Features using Kernels: Landmarks

- Suppose we choose a small set of  $L$  “landmark” inputs  $\mathbf{z}_1, \mathbf{z}_2, \dots, \mathbf{z}_L$  in the training data



$$\psi(\mathbf{x}_n) = [k(\mathbf{z}_1, \mathbf{x}_n), k(\mathbf{z}_2, \mathbf{x}_n), k(\mathbf{z}_3, \mathbf{x}_n)] \in \mathbb{R}^3$$

- For each input  $\mathbf{x}_n$ , using a kernel  $k$ , define an  $L$ -dimensional feature vector as follows

$$\psi(\mathbf{x}_n) = [k(\mathbf{z}_1, \mathbf{x}_n), k(\mathbf{z}_2, \mathbf{x}_n), \dots, k(\mathbf{z}_L, \mathbf{x}_n)] \in \mathbb{R}^L$$

- Can now apply a linear model on  $\psi$  representation ( $L$ -dimensional now) of the inputs
- This will be fast both at training as well as test time if  $L$  is small
- No need to kernelize the linear model while still reaping the benefits of kernels 😊



# Extracting Feat. using Kernels: Random Features

- Many kernel functions\* can be written as

$$k(\mathbf{x}_n, \mathbf{x}_m) = \phi(\mathbf{x}_n)^\top \phi(\mathbf{x}_m) = \mathbb{E}_{\mathbf{w} \sim p(\mathbf{w})} [t_{\mathbf{w}}(\mathbf{x}_n) t_{\mathbf{w}}(\mathbf{x}_m)]$$

.. where  $t_{\mathbf{w}}(\cdot)$  is a function with params  $\mathbf{w} \in \mathbb{R}^L$  with  $\mathbf{w}$  drawn from some distr.  $p(\mathbf{w})$

- Example: For the RBF kernel,  $t_{\mathbf{w}}(\cdot)$  is cosine func. and  $p(\mathbf{w})$  is zero mean Gaussian


$$k(\mathbf{x}_n, \mathbf{x}_m) = \mathbb{E}_{\mathbf{w} \sim p(\mathbf{w})} [\cos(\mathbf{w}^\top \mathbf{x}_n) \cos(\mathbf{w}^\top \mathbf{x}_m)]$$

- Given  $\mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_L$  from  $p(\mathbf{w})$ , using Monte-Carlo approx. of above expectation

$$k(\mathbf{x}_n, \mathbf{x}_m) \approx \frac{1}{L} \sum_{\ell=1}^L \cos(\mathbf{w}_\ell^\top \mathbf{x}_n) \cos(\mathbf{w}_\ell^\top \mathbf{x}_m) = \psi(\mathbf{x}_n)^\top \psi(\mathbf{x}_m)$$

.. where  $\psi(\mathbf{x}_n) = \frac{1}{\sqrt{L}} [\cos(\mathbf{w}_1^\top \mathbf{x}_n), \dots, \cos(\mathbf{w}_L^\top \mathbf{x}_n)]$  is an  $L$ -dim vector

- Can apply a linear model on this  $L$ -dim rep. of the inputs (no need to kernelize)



# Learning with Kernels: Some Aspects

- Storage/computational efficiency can be a bottleneck when using kernels
- During training, need to compute and store the  $N \times N$  kernel matrix  $\mathbf{K}$  in memory
- Need to store training data (or at least support vectors in case of SVMs) at test time
- Test time can be slow:  $O(N)$  cost to compute a quantity like  $\sum_{n=1}^N \alpha_n k(\mathbf{x}_n, \mathbf{x}_*)$
- Approaches like landmark and random features can be used to speed up
- Choice of the right kernel is also very important
- Some kernels (e.g., RBF) work well for many problems but hyperparameters of the kernel function may need to be tuned via cross-validation
- Quite a bit of research on learning the right kernel from data
  - Learning a combination of multiple kernels ([Multiple Kernel Learning](#))
  - [Bayesian kernel methods](#) (e.g., [Gaussian Processes](#)) can learn the kernel hyperparameters from data (thus can be seen as learning the kernel)
  - Deep Learning can also be seen as learning the kernel from data (more on this later)

Also, a lot of recent work on connections between kernel methods and deep learning



# Coming up next

- Unsupervised learning

